


SST/macro 7.2: User's Manual

Sandia National Labs
Livermore, CA

July 24, 2017



Contents

1	Introduction	6
1.1	Overview	6
1.2	Preview of Things to Come	7
1.3	What To Expect In The User's Manual	8
2	Building and Running SST/macro	9
2.1	Build and Installation of SST/macro	9
2.1.1	Downloading	9
2.1.2	Dependencies	9
2.1.3	Configuration and Building	10
2.1.4	Post-Build	11
2.1.5	GNU pthread for user-space threading	12
2.1.6	Building Boost (no longer required)	12
2.1.7	Known Issues	13
2.2	Building DUMPI	13
2.2.1	Known Issues	13
2.3	Building with OTF2 (Beta)	14
2.4	Building Clang source-to-source support (Beta)	14
2.4.1	Building Clang libTooling	14
2.4.2	Building SST/macro with Clang	15
2.5	Running an Application	16
2.5.1	SST Python Scripts	16
2.5.2	Building Skeleton Applications	16
2.5.3	Makefiles	17

2.5.4	Command-line arguments	17
2.6	Parallel Simulations in Standalone Mode	17
2.6.1	Distributed Memory Parallel	18
2.6.2	Shared Memory Parallel	18
2.6.3	Warnings for Parallel Simulation	18
2.7	Debug Output	19
3	Basic Tutorials	20
3.1	SST/macro Parameter files	20
3.1.1	Naming Conventions	21
3.1.2	Parameter Namespace Rules	21
3.1.3	Initial Example	22
3.2	Abstract Machine Models	23
3.2.1	Common Parameters	23
3.2.2	AMM1	24
3.2.3	AMM2	25
3.2.4	AMM3	26
3.3	Network Topologies and Routing	27
3.3.1	Topology	27
3.3.2	Routing	29
3.4	Network Model	30
3.4.1	Analytic Models: MACRELS	30
3.4.2	Packet Models: PISCES	31
3.4.3	Flow	33
3.5	Basic MPI Program	34
3.6	Launching, Allocation, and Indexing	35
3.6.1	Launch Commands	35
3.6.2	Allocation Schemes	35
3.7	Discrete Event Simulation	38
3.8	Using DUMPI	40
3.8.1	Building DUMPI	40
3.8.2	Trace Collection	40
3.8.3	Trace Replay	42

3.9	Using Score-P and OTF2 (Beta)	43
3.9.1	Trace Collection	44
3.9.2	Trace Replay	44
3.10	Call Graph Visualization	45
3.11	Spyplot Diagrams	47
3.12	Fixed-Time Quanta Charts	48
3.13	Network Statistics	49
3.13.1	Message Size Histogram	49
3.13.2	Congestion Delay Histogram	49
3.13.3	Congestion Spyplot and Multi-stats	50
4	Topologies	53
4.1	Topology Query Utility	53
4.2	Torus	54
4.3	Hypercube	55
4.3.1	Allocation and indexing	56
4.3.2	Routing	57
4.4	Fat Tree	57
4.4.1	Allocation and indexing	60
4.4.2	Routing	62
4.5	Dragonfly	62
4.5.1	Allocation and indexing	62
4.5.2	Routing	63
5	External Applications and Skeletonization	67
5.1	Basic Application porting	67
5.1.1	Loading external skeletons with the standalone core	67
5.1.2	Loading external skeletons with the integrated core	68
5.2	Auto-skeletonization with Clang (Beta)	69
5.2.1	Redirecting Main	69
5.2.2	Memory Allocations	69
5.2.3	Computation	69
5.2.4	Special Pragmas	70
5.2.5	Skeletonization Issues	70
5.3	Process Encapsulation	71

6	Clang Source-to-Source Auto-Skeletonization via Pragmas	72
6.1	Pragma Overview	72
6.1.1	Compiler workflow	72
6.1.2	Compiler Environment Variables	72
6.2	Basic Replacement Pragmas	74
6.2.1	pragma sst delete: no arguments	74
6.2.2	pragma sst replace [to_replace:string] [new_text:C++ expression]	74
6.2.3	pragma sst init [new_value:string]	75
6.2.4	pragma sst return [new_value:C++ expression]	75
6.2.5	pragma sst keep	75
6.2.6	pragma sst keep_if [condition:C++ bool expression]	75
6.2.7	pragma sst empty	75
6.2.8	pragma sst branch_predict [condition:C++ expression]	76
6.3	Memory Allocation Pragmas	76
6.3.1	pragma sst malloc	76
6.3.2	pragma sst new	76
6.4	Data-Driven Type Pragmas	76
6.4.1	pragma sst null_variable	76
6.4.2	pragma sst null_type [type alias] [list allowed functions]	77
6.5	Compute Pragmas	77
6.5.1	pragma sst compute and pragma omp parallel	77
6.5.2	pragma sst loop_count [integer: C++ expression]	78
6.5.3	pragma sst branch_predict [float: C++ expression]	78
7	Issues and Limitations	79
7.1	Polling in applications	79
7.2	Fortran	79

8	Detailed Parameter Listings	80
8.1	Global namespace	81
8.2	Namespace “topology”	82
8.3	Namespace “node”	83
8.3.1	Namespace “node.nic”	83
8.3.2	Namespace “node.memory”	88
8.3.3	Namespace “node.os”	88
8.3.4	Namespace “node.proc”	89
8.4	Namespace “netlink”	89
8.4.1	Namespace “netlink.injection”	90
8.4.2	Namespace “netlink.ejection”	91
8.5	Namespace “mpi”	91
8.5.1	Namespace “mpi.queue”	92
8.6	Namespace “switch”	92
8.6.1	Namespace “switch.router”	92
8.6.2	Namespace “switch.output_buffer”	92
8.6.3	Namespace “switch.xbar”	94
8.6.4	Namespace “switch.link”	95
8.6.5	Namespace “switch.ejection”	95
8.7	Namespace “appN”	96

Chapter 1

Introduction

1.1 Overview

The SST/macro software package provides a simulator for large-scale parallel computer architectures. It permits the coarse-grained study of distributed-memory applications. The simulator is driven from either a trace file or skeleton application. The simulator architecture is modular, allowing it to easily be extended with additional network models, trace file formats, software services, and processor models.

Simulation can be broadly categorized as either off-line or on-line. Off-line simulators typically first run a full parallel application on a real machine, recording certain communication and computation events to a simulation trace. This event trace can then be replayed post-mortem in the simulator. Most common are MPI traces which record all MPI events, and SST/macro provides the DUMPI utility (3.8) for collecting and replaying MPI traces. Trace extrapolation can extend the usefulness of off-line simulation by estimating large or untraceable system scales without having to collect a trace, but it is limited.

We turn to on-line simulation when the hardware or applications parameters need to change. On-line simulators instead run real application code, allowing native C/C++/Fortran to be compiled directly into the simulator. SST/macro intercepts certain function calls, estimating how much time passes rather than actually executing the function. In MPI programs, for example, calls to MPI_Send are linked to the simulator instead of passing to the real MPI library. If desired, SST/macro can actually be a full MPI *emulator*, delivering messages between ranks and replicating the behavior of a full MPI implementation.

Although SST/macro supports both on-line and off-line modes, on-line simulation is encouraged because event traces are much less flexible, containing a fixed sequence of events. Application inputs and number of nodes cannot be changed. Without a flexible control flow, it also cannot simulate dynamic behavior like load balancing or faults. On-line simulation can explore a much broader problem space since they evolve directly in the simulator.

For large, system-level experiments with thousands of network endpoints, high-accuracy cycle-accurate simulation is not possible, or at least not convenient. Simulation requires coarse-grained approximations to be practical. SST/macro is therefore designed for specific cost/accuracy tradeoffs. It should still capture complex cause/effect behavior in applications and hardware, but be efficient enough to simulate at the system-level. For speeding up simulator execution, we encourage *skeletonization*, discussed further in Chapter 5. A high-quality skeleton is an application model that reproduces certain characteristics with only limited computation. We also encourage uncertainty quantification (UQ) for validating simulator results. Skeletonization and UQ are the two main elements in the “canonical” SST/macro workflow (Figure 1.1).



Figure 1.1: SST/macro workflow.

Because of its popularity, MPI is one of our main priorities in providing programming model support. Some MPI-3 functions and MPI one-sided functions are not implemented. This will lead to compile errors with an obvious “not implement” compiler message.

1.2 Preview of Things to Come

Suppose you have the basic MPI application below that executes a simple send/rcv operation. One could use `mpicc` or `mpic++` to compile and run as an actual MPI program. This requires spawning all the processes and running them in parallel. Suppose, however, you wanted to simulate an entire MPI job launch within a single process. This might prove very useful for debugging since you could just run GDB or Valgrind on a single process. It might take a while, but for small runs (16 ranks or so) you could debug right on your laptop the same you do for a serial program.


```

1  int size = atoi(argv[1]);
2  if (rank == 0){
3      int partner = 1;
4      MPI_Send(buffer, size, MPI_INT, partner, tag, MPI_COMM_WORLD);
5  } else {
6      int partner = 0;
7      MPI_Recv(buffer, size, MPI_INT, partner, tag, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
8  }
9  MPI_Barrier(MPI_COMM_WORLD);
10
11 if (rank == 0){
12     printf("Rank 0 finished at t=%8.4f ms\n", MPI_Wtime()*1e3);
13 }

```

This is exactly the functionality that SST/macro provides. Instead of using `mpic++`, you compile the code with `sst++`. This modifies your code and intercepts MPI calls, running them through the simulator instead of an actual MPI implementation. Your code will execute and run exactly the same. Your application won't even know the difference. The only major hiccup is that you now need a parameter file with information like:

```

node {
  appl {
    name = send_recv
    launch_cmd = aprun -n 2
    argv = 20
  }
}

```

Rather than launching your code using `mpirun` or similar, you put all your command line parameters into a `parameters.ini` file and run:

```
shell> sstmac -f parameters.ini
```

The simulator then executes your application exactly as if you had been a real system and run:

```
shell> aprun -n 2 ./send_recv 20
```

Things get more complicated when you bring skeletonization into play. The use case above is *emulation*, exactly reproducing MPI functionality. In *skeletonization* or *simulation*, SST/macro will mimic as closely as possible the original application, but avoids as much computation and as much memory allocation as possible. This allows to pack in as many simulated MPI ranks as possible into your single `sstmac` process.

1.3 What To Expect In The User's Manual

This user's manual is mainly designed for those who wish to perform experiments with *new* applications using *existing* hardware models. This has been the dominant use case and we therefore classify those doing application experiments as "users" and those making new hardware models "developers." Getting applications to run in SST/macro should be very straightforward and requires no knowledge of simulator internal code. Making new hardware models is much more in depth and requires learning some basics of core simulator code. Those interested in making new hardware models should consult the developer's manual in the top-level source directory.

Chapter 2

Building and Running SST/macro

2.1 Build and Installation of SST/macro

2.1.1 Downloading

SST/macro is available at <https://github.com/sstsimulator/sst-macro>. You can download the git repository directly:

```
shell> git clone https://github.com/sstsimulator/sst-macro.git
```

or for ssh

```
shell> git clone ssh://git@github.com/sstsimulator/sst-macro.git
```

or you can download a release tarball from <https://github.com/sstsimulator/sst-macro/releases>.

2.1.2 Dependencies

The most critical change is that C++11 is now a strict prerequisite. Workarounds had previously existed for older compilers. These are no longer supported. The following are dependencies for SST/macro.

- (optional) Git is needed in order to clone the source code repository, but you can also download a tar (Section 2.1.1).
- (Mac) You need the GNU pth library. Downloadable from MacPorts (see Section 2.1.5)
- Autoconf: 2.68 or later
- Automake: 1.11.1 or later
- Libtool: 2.4 or later
- A C/C++ compiler is required with C++11 support. gcc >=4.8 and clang >= 3.7 are known to work.
- (optional) OTF2: 2.0 or later for OTF2 trace replay.

- (optional) Doxygen and Graphviz are needed to build the documentation.
- (optional) KCacheGrind or QCacheGrind to display call graphs
- (optional) Clang development libraries to enable SST source-to-source compiler

2.1.3 Configuration and Building

SST/macro is an SST element library, proving a set of simulation components that run on the main SST core. The SST core provides the parallel discrete event simulation manager that manages time synchronization and sending events in serial, MPI parallel, multi-threaded, or MPI + threaded mode. The core does not provide any simulation components like node models, interconnect models, MPI trace readers, etc. The actual simulation models are contained in the element library.

The SST core is a standalone executable that dynamically loads shared object files containing the element libraries. For many element libraries, a Python input file is created that builds and connects the various simulation components. For maximum flexibility, this will become the preferred mode. However, SST/macro has historically had a text-file input `parameters.ini` that configures the simulation. To preserve that mode for existing users, a wrapper Python script is provided that processes SST/macro input files. SST/macro can also be compiled in standalone mode that uses its own simulation core.

The workflow for installing and running on the main SST core is:

- Build and install SST core
- Build and install the SST/macro element library `libmacro.so`
- Make sure paths are properly configured for `libmacro.so` to be visible to the SST core (`SST_LIB_PATH`)
- Run the `pysstmac` wrapper Python script that runs SST/macro-specific parameters OR
- Write a custom Python script

The workflow for installing and running on the standalone SST/macro core:

- Build and install SST/macro standalone to generate `sstmac` executable
- Run `sstmac` with `*.ini` parameter files

Build SST core

The recommended mode for maximum flexibility is to run using the SST core downloadable from <http://sst-simulator.org/SSTPages/SSTMainDownloads>. Building and installing sets up the discrete event simulation core required for all SST elements. SST core no longer has Boost dependencies! Directions for building Boost (if desired) are still below in 2.1.6

Build SST/macro element library

If using the repo (not a release tarball), go to the top-level source directory and run:

```
top-level-src> ./bootstrap.sh
```

This sets up the configure script. For builds, we recommend building outside the source tree:

```
sst-macro> mkdir build
sst-macro> cd build
sst-macro/build> ../configure --prefix=$PATH_TO_INSTALL --with-sst-core=$PATH_TO_SST_CORE CC=$MPICC CXX=$MPICXX
```

`PATH_TO_SST_CORE` should be the prefix install directory used when installing the core. The MPI compilers should be the same compilers used for building Boost and SST core.

SST/macro can still be built in standalone mode for a select set of features that have not been fully migrated to the SST core. The installation and running are the same for both modes - simply remove the `--with-sst-core` parameter. A complete list of options for standalone building can be seen by running `../configure --help`. Some common options include:

- `-(dis|en)able-graphviz` : Enables the collection of simulated call graphs, which can be viewed with graphviz. Enabled by default. Disable if not using Boost or C++11. Ordered maps can be used as a replacement, but with lower performance.
- `-(dis|en)able-regex` : Regular expressions can be used to proofread input files, but this requires either Boost or C++11. Enabled by default. Disable if not using Boost or C++11.
- `-(dis|en)able-custom-new` : Memory is allocated in larger chunks in the simulator, which can speed up large simulations.
- `-(dis|en)able-multithread` : This configures for thread-level parallelism for (hopefully) faster simulation
- `-(dis|en)able-otf2[=location]`: Enable OTF2 trace replay, requires a path to OTF2 installation.
- `-with-clang[=location]`: Enable Clang source-to-source tools by pointing to Clang development libraries

Once configuration has completed, printing a summary of the things it found, simply type `make`.

2.1.4 Post-Build

If the build did not succeed open an issue on the github page at <https://github.com/sstsimulator/sst-macro/issues> or contact SST/macro support for help (sst-macro-help@sandia.gov).

If the build was successful, it is recommended to run the range of tests to make sure nothing went wrong. To do this, and also install SST/macro to the install path specified during installation, run the following commands:

```
sst-macro/build> make check
sst-macro/build> make install
sst-macro/build> make installcheck
```

`Make check` runs all the tests we use for development, which checks all functionality of the simulator. `Make installcheck` compiles some of the skeletons that come with SST/macro, linking against the installation.

Important: Applications and other code linking to SST/macro use Makefiles that use the sst++/sstcc compiler wrappers that are installed there for convenience to figure out where headers and libraries are. When making your skeletons and components, make sure your path is properly configured.

2.1.5 GNU pthread for user-space threading

By default, Linux usually ships with `ucontext` which enables user-space threading. Mac OS X does not support `ucontext` and requires an extra library be installed (GNU pthread). GNU pthread is easy to download and install from source. Even easier is MacPorts.

```
shell> sudo port install pthread
```

MacPorts installed all libraries to `/opt/local`. When configuring, simply add `--with-pthread=$PATH_TO_PTH` as an argument. For MacPorts installation, this means configuring SST/macro with:

```
../configure --with-pthread=/opt/local
```

2.1.6 Building Boost (no longer required)

We recommend two files: `user-config.jam` to configure the Boost compiler flags and a `runme.sh` that bootstraps, compiles, and installs the prerequisite Boost libraries. For GCC, the `user-config.jam` should go in the top-level home directory and the file should contain the line:

```
using gcc : : $PATH_TO_MPIC++ : <compileflags>-std=c++1y ;
```

For Clang on Mac, use:

```
using clang : : $PATH_TO_CLANG_MPIC++ : <compileflags>-std=c++1y <linkflags>-stdlib=libc++
```

We recommend Boost 1.59. Other Boost versions should work as well, but this seems the most stable. In the top-level Boost directory, make a script `runme.sh` that contains:

```
./bootstrap.sh \
  --with-libraries=program_options,serialization,filesystem \
  --with-toolset=gcc

./b2 ---prefix=$INSTALL
./b2 ---layout=tagged ---prefix=$INSTALL
```

The toolset can be changed from gcc to clang, as needed. For maximum safety, Boost should install both “tagged” versions of libraries and un-tagged versions. Once Boost is installed with these options, configuration and installation of SST core should be straightforward following documentation in the core library.

2.1.7 Known Issues

- Any build or runtime problems should be reported to sst-macro-help@sandia.gov. We try to respond as quickly as possible.
- make -j: When doing a parallel compile dependency problems can occur. There are a lot of inter-related libraries and files. Sometimes the Makefile dependency tracker gets ahead of itself and you will get errors about missing libraries and header files. If this occurs, restart the compilation. If the error vanishes, it was a parallel dependency problem. If the error persists, then it's a real bug.
- GNU pth: For Mac, make sure to follow directions in 2.1.5 to ensure pth is correctly installed.
- Ubuntu: The Ubuntu linker performs too much optimization on dynamically linked executables. Some call it a feature. I call it a bug. In the process it throws away symbols it actually needs later. The build system should automatically fix Ubuntu flags. If still having issues, make sure that '-Wl,-no-as-needed' is given in LDFLAGS.

The problem occurs when the executable depends on libA which depends on libB. The executable has no direct dependence on any symbols in libB. Even if you add -lB to the LDFLAGS or LDADD variables, the linker ignores them and throws the library out. It takes a dirty hack to force the linkage. If there are issues, contact the developers at sst-macro-help@sandia.gov and report the problem.

2.2 Building DUMPI

By default, DUMPI is configured and built along with SST/macro with support for reading and parsing DUMPI traces, known as libundumpi. DUMPI binaries and libraries are also installed along with everything for SST/macro during make install. DUMPI can be used as its own library within the SST/macro source tree by changing to `sst-macro/sst-dumpi`, where you can change its configuration options.

DUMPI can also be used as stand-alone tool (*e.g.* for simplicity if you're only tracing). To get DUMPI by itself, either copy the `sstmacro/dumpi` directory somewhere else or visit <https://github.com/sstsimulator/sst-dumpi> and follow similar instructions for obtaining SST/macro.

To see a list of configuration options for DUMPI, run `./configure --help`. If you're trying to configure DUMPI for trace collection, use `--enable-libdumpi`. Your build process might look like this (if you're building in a separate directory from the dumpi source tree) :

```
sst-dumpi/build> ../configure --prefix=/path-to-install --enable-libdumpi
sst-dumpi/build> make
sst-dumpi/build> sudo make install
```

2.2.1 Known Issues

- When compiling on platforms with compiler/linker wrappers, *e.g.* ftn (Fortran) and CC (C++) compilers at NERSC, the libtool configuration can get corrupted. The linker flags automatically added by the wrapper produce bad values for the `predeps/postdeps` variable in the libtool script in the top level source folder. When this occurs, the (unfortunately) easiest way to fix this is to manually modify the libtool script. Search for `predeps/postdeps` and set the values to empty. This will clear all the erroneous linker flags. The compilation/linkage should still work since all necessary flags are set by the wrappers.

2.3 Building with OTF2 (Beta)

OTF2 is a general purpose trace format with specialized callbacks for the MPI API. OTF2 traces are generated by programs compiled with Score-P compiler wrappers. SST/macro 7.0 supports OTF2 trace replay when configured with

```
./configure --enable-otf2=<OTF2-root>
```

where the OTF2 root is the installation prefix for a valid OTF2 build. OTF2 can be obtained from the Score-P project at <http://www.vi-hps.org/projects/score-p>. Detailed build and usage instructions can be found on the website.

2.4 Building Clang source-to-source support (Beta)

To enable Clang source-to-source support it is not sufficient to have a Clang compiler. You have to install a special libTooling library for Clang.

2.4.1 Building Clang libTooling

The Easy Way: Mac OS X

Using MacPorts on OS X, it is trivial to obtain a Clang installation that includes libTooling:

```
port install clang-devel
```

MacPorts will place the Clang compilers in `/opt/local/bin`. Enable the devel version of Clang with:

```
port select --set clang mp-clang-devel
```

The Clang libraries will be placed into `/opt/local/libexec/llvm-devel/lib`, so the appropriate option to the sst-macro configure script is `--with-clang=/opt/local/libexec/llvm-devel`.

The Hard Way

For operating systems other than OS X, building Clang support has a few steps (and takes quite a while to build), but is straightforward. Instead of having an all-in-one tarball, you will have to download 2 different components. You can install more if you want build libc++, but these are not required. Obtain the following from <http://releases.llvm.org/download.html>.

- LLVM source code
- Clang source code
- (optional, not recommended unless needed) libc++ source code
- (optional, not recommended unless needed) libc++abi source code
- (optional, not recommended) compiler-rt source code

- (optional, not recommended) OpenMP source code

Setting up the folders can be done automatically using the `setup-clang` script in `bin/tools` folder in `sst-macro`. Put all of downloaded tarballs in a folder, e.g. `clang-llvm`. Then run `setup-clang` in the directory. It will automatically place files where LLVM needs them. LLVM is the “driver” for the entire build. Everything else is a subcomponent. The setup script places each tarball in the following subfolders of the main LLVM folder

- `tools/clang`
- `projects/compiler-rt`
- `projects/libc++`
- `projects/libc++abi`
- `projects/openmp`

Only Clang is a strict dependency. Using CMake (assuming you are in a build subdirectory of the LLVM tree), you would run the script below to configure. You should not longer need to use Clang to build Clang, but for the most stable results you should a pre-existing Clang compiler to build the Clang development libraries.

```
cmake ../llvm \
-DMAKE_CXX_COMPILER=clang++ \
-DMAKE_C_COMPILER=clang \
-DMAKE_CXX_FLAGS="-O3" \
-DMAKE_C_FLAGS="-O3" \
-DMAKE_INSTALL_PREFIX=$install
```

To build a complete LLVM/Clang (again, not required unless you absolutely need a new libc++), run:

```
cmake ../llvm \
-DMAKE_CXX_COMPILER=clang++ \
-DMAKE_C_COMPILER=clang \
-DMAKE_CXX_FLAGS="-O3" \
-DMAKE_C_FLAGS="-O3" \
-DLLVM_ENABLE_LIBCXX=ON \
-DLLVM_TOOL_COMPILER_RT_BUILD=ON \
-DLLVM_TOOL_LIBCXXABI_BUILD=ON \
-DLLVM_TOOL_LIBCXX_BUILD=ON \
-DMAKE_INSTALL_PREFIX=$install
```

On some systems, linking Clang might blow out your memory. If that is the case, you have to set `LD=ld.gold` for the linker. Run `make install`. The libTooling library will now be available at the `$install` location.

2.4.2 Building SST/macro with Clang

Now that clang is installed, you only need to add the configure flag `--with-clang` pointing it to the install location from above. You must use the same Clang compiler to build SST that you used to build libTooling.

```
../configure CXX=clang++ CC=clang --with-clang=$install
```

Clang source-to-source support will now be built into the `sst++` compiler. If Clang development libraries are available in the default system path (as is often the case with LLVM models, e.g. `module load llvm`), then you can just put `--with-clang`.

2.5 Running an Application

2.5.1 SST Python Scripts

Full details on building SST Python scripts can be found at <http://sst-simulator.org>. To preserve the old parameter format in the near-term, SST/macro provides the `pysstmac` script:

```
export SST_LIB_PATH=$SST_LIB_PATH:$SSTMAC_PREFIX/lib
options="$@"
$SST_PREFIX/bin/sst $SSTMAC_PREFIX/include/python/default.py --model-options="$options"
```

The script configures the necessary paths and then launches with a Python script `default.py`. Interested users can look at the details of the Python file to understand how SST/macro converts parameter files into a Python config graph compatible with SST core. Assuming the path is configured properly, users can run

```
shell>pysstmac -f parameters.ini
```

with a properly formatted parameter file. If running in standalone mode, the command would be similar (but different).

```
shell>sstmac -f parameters.ini
```

since there is no Python setup involved.

2.5.2 Building Skeleton Applications

To demonstrate how an external skeleton application is run in SST/macro, we'll use a very simple send-recv program located in `skeletons/sendrecv`. We will take a closer look at the actual code in Section 3.5. After SST/macro has been installed and your PATH variable set correctly, for standalone core users can run:

```
sst-macro> cd skeletons/sendrecv
sst-macro/skeletons/sendrecv> make
sst-macro/skeleton/sendrecv> ./runsstmac -f parameters.ini
```

You should see some output that tells you 1) the estimated total (simulated) runtime of the simulation, and 2) the wall-time that it took for the simulation to run. Both of these numbers should be small since it's a trivial program.

This is how simulations generally work in SST/macro: you build skeleton code and link it with the simulator to produce a binary. Then you run that binary and pass it a parameter file which describes the machine model to use. For running on the main SST core, a few extra flags are required. Rather than generating a standalone executable, the compiler wrapper generates a shared library. Users can always write their own Python scripts, which will be required for more advanced usage. However, users can also just use the `pysstmac` script.

```
sst-macro/skeletons/sendrecv> pysstmac librunsstmac.so -f parameters.ini
```

Any extra shared libraries can be given as the first few parameters and these will automatically be imported.

2.5.3 Makefiles

We recommend structuring the Makefile for your project like the one seen in `skeletons/sendrecv/Makefile` :

```
TARGET := runsstmac
SRC := $(shell ls *.c)

CXX :=      $(PATH_TO_SST)/bin/sst++
CC :=      $(PATH_TO_SST)/bin/sstcc
CXXFLAGS := ...
CPPFLAGS := ...
LIBDIR :=  ...
PREFIX :=  ...
LDFLAGS := -Wl,-rpath,$(PREFIX)/lib
...
```

The SST compiler wrappers `sst++` and `sstcc` automatically configure and map the code for simulation.

2.5.4 Command-line arguments

There are only a few basic command-line arguments you'll ever need to use with SST/macro, listed below

- `-h/-help`: Print some typical help info
- `-f [parameter file]`: The parameter file to use for the simulation. This can be relative to the current directory, an absolute path, or the name of a pre-set file that is in `sstmacro/configurations` (which installs to `/path-to-install/include/configurations`, and gets searched along with current directory).
- `-dumpi`: If you are in a folder with all the DUMPI traces, you can invoke the main `sstmac` executable with this option. This replays the trace in a special debug mode for quickly validating the correctness of a trace.
- `-d [debug flags]`: A list of debug flags to activate as a comma-separated list (no spaces) - see Section 2.7
- `-p [parameter]=[value]`: Setting a parameter value (overrides what is in the parameter file)
- `-t [value]`: Stop the simulation at simulated time [value]
- `-c`: If multithreaded, give a comma-separated list (no spaces) of the core affinities to use - see Section 2.6.2

2.6 Parallel Simulations in Standalone Mode

SST/macro supports running parallel discrete event simulation (PDES) in distributed memory (MPI), threaded shared memory (pthreads) and hybrid (MPI+pthreads) modes. Running these in standalone mode will be discouraged as parallel simulations should use the unified SST core. However, near-term, hybrid modes and other optimizations are not fully supported in the unified SST core. They standalone core may still be required for certain cases.

2.6.1 Distributed Memory Parallel

Configure will automatically check for MPI. Your configure should look something like:

```
sst-macro/build> ../configure CXX=mpicxx CC=mpicc ...
```

With the above options, you can just compile and go. SST/macro is run exactly like the serial version, but is spawned like any other MPI parallel program. Use your favorite MPI launcher to run, e.g. for OpenMPI

```
mysim> mpirun -n 4 sstmac -f parameters.ini
```

or for MPICH

```
mysim> mpiexec -n 4 sstmac -f parameters.ini
```

Even if you compile for MPI parallelism, the code can still be run in serial with the same configuration options. SST/macro will notice the total number of ranks is 1 and ignore any parallel options. When launched with multiple MPI ranks, SST/macro will automatically figure out how many partitions (MPI processes) you are using, partition the network topology into contiguous blocks, and start running in parallel.

2.6.2 Shared Memory Parallel

In order to run shared memory parallel, you must configure the simulator with the `--enable-multithread` flag. Partitioning for threads is currently always done using block partitioning and there is no need to set an input parameter. Including the integer parameter `sst_nthread` specifies the number of threads to be used (per rank in MPI+threads mode) in the simulation. The following configuration options may provide better threaded performance.

- `--enable-spinlock` replaces pthread mutexes with spinlocks. Higher performance and recommended when supported.
- `--enable-cpu-affinity` causes SST/macro to pin threads to specific cpu cores. When enabled, SST/macro will require the `cpu_affinity` parameter, which is a comma separated list of cpu affinities for each MPI task on a node. SST/macro will sequentially pin each thread spawned by a task to the next next higher core number. For example, with two MPI tasks per node and four threads per MPI task, `cpu_affinity = 0,4` will result in MPI tasks pinned to cores 0 and 4, with pthreads pinned to cores 1-3 and 5-7. For a threaded only simulation `cpu_affinity = 4` would pin the main process to core 4 and any threads to cores 5 and up. The affinities can also be specified on the command line using the `-c` option. Job launchers may in some cases provide duplicate functionality and either method can be used.

2.6.3 Warnings for Parallel Simulation

- If the number of simulated processes specified by e.g. `aprun -n 100` does not match the number of nodes in the topology (i.e. you are not space-filling the whole simulated machine), parallel performance will suffer. SST/macro partitions nodes, not MPI ranks.

Parallel simulation speedups are likely to be modest for small runs. Speeds are best with serious congestion or heavy interconnect traffic. Weak scaling is usually achievable with 100-500 simulated MPI ranks per logical process. Even without speedup, parallel simulation can certainly be useful in overcoming memory constraints, expanding the maximum memory footprint.

2.7 Debug Output

SST/macro defines a set of debug flags that can be specified in the parameter file to control debug output printed by the simulator. To list the set of all valid flags with documentation, the user can run

```
bin> ./sstmac --debug-flags
```

which will output something like

```
mpi
  print all the basic operations that occur on each rank – only API calls are
  logged, not any implementation details
mpi_check
  validation flag that performs various sanity checks to ensure MPI application
  runs and terminates cleanly
mpi_collective
  print information about MPI collective calls as well as implementation details
mpi_pt2pt
  print information about MPI point-to-point calls as well as implementation
  details
....
```

The most important flag for validating simulations is the `mpi_check` flag, which causes special sanity checks and a final validation check to ensure the simulation has finished cleanly. Some of the debug flags can generate information overload and will only be useful to a serious developer, rather than a user.

To turn on debug output, add the following to the input file

```
debug = mpi  mpi_check
```

listing all flags you want separated by spaces. Note: this is a major shift from the previous (and really tedious, unfriendly) debug system of past versions. The new system allows much finer-grained, simpler printing of debug output. Additionally, it allows new debug flags to very easily defined. More info on declaring new debug flags in your own code can be found in the developer's reference.

Chapter 3

Basic Tutorials

3.1 SST/macro Parameter files

A minimal parameter file setting up a 2D-torus topology is shown below. A detailed listing of parameter namespaces and keywords is given in Section 8. The preferred input files now use namespaces. However, for consistency with previous versions, we also show the deprecated parameters.

```
amm_model = amml
congestion_model = LogP
node {
  #run a single mpi test
  appl {
    indexing = block
    allocation = first_available
    launch_cmd = aprun -n8 -N1
    name = sstmac_mpi_testall
    argv =
    sendrecv_message_size = 128
  }
  ncores = 1
  memory {
    model = simple
    bandwidth = 1GB/s
    latency = 10ns
  }
  proc {
    frequency = 1GHz
  }
  nic {
    injection {
      bandwidth = 1GB/s
      latency = 1us
    }
    model = simple
  }
}
switch {
  ejection {
    bandwidth = 1GB/s
    latency = 100ns
  }
  link {
    bandwidth = 1.0GB/s
  }
}
```

```

    latency = 100ns
  }
}

topology {
  name = torus
  geometry = 4,4
}

```

The input file follows a basic syntax of **parameter = value**. Parameter names follow C++ variable rules (letters, numbers, underscore) while parameter values can contain spaces. Trailing and leading whitespaces are stripped from parameters. Comments can be included on lines starting with `#`.

3.1.1 Naming Conventions

There are some patterns in the naming of SST parameters. The most critical is “model” versus “name”. Consider:

```

1 switch {
2   model = pisces
3   router {
4     name = minimal
5   }
6 }
7

```

The distinction here is subtle, but important. When a parameter has the name “model” it means there are important approximations being made. The model does not necessarily reproduce a real switch or node or NIC. It’s not a choice of algorithm. It’s the choice of an approximate hardware model. In contrast, when selecting an algorithm that is trying to reproduce “exact” behavior, the parameter is “name.” This indicates the name of an algorithm, not an approximate model.

3.1.2 Parameter Namespace Rules

Periods denote nesting of parameter namespaces. The parameter `node.memory.model` will be nested in namespace `memory` inside namespace `node`. If inside a given namespace, SST/macro looks within that namespace first. If unable to find it, the input parser moves through the namespace nesting to find the value. This exactly follows C++ namespace rules with one important exception. The “global” namespace is reserved for special keywords that get translated. Keywords in the global namespace are not visible to any components.

For example, consider the following:

```

node.model = simple
node.memory.model = pisces

```

If I am building the node’s memory system, the initialization will look for `memory::model` inside namespace `node` first, returning the value `pisces`.

The preferred syntax from 6.1 more closely resembles C++ namespace declarations. Namespaces can be scoped using brackets `{}`:

```

node {
  model = simple
  memory {
    model = simple
    bandwidth = 1GB/s
    latency = 10ns
  }
}

```

Any line containing a single string with an opening `{` starts a new namespace. A line containing only a closing `}` ends the innermost namespace. The syntax is not as flexible as C++ since the opening `{` must appear on the same line as the namespace and the closing `}` must be on a line of its own. A detailed listing of parameter namespaces and keywords is given in Section 8.

3.1.3 Initial Example

Continuing with the example above, we see the input file is broken into namespace sections. First, application launch parameters for each node must be chosen determining what application will launch, how nodes will be allocated, how ranks will be indexed, and finally what application will be run. Additionally, you must specify how many processes to launch and how many to spawn per node. We currently recommend using `aprun` syntax (the launcher for Cray machines), although support is being added for other process management systems. `SST/macro` can simulate command line parameters by giving a value for `node.app1.argv`.

A network must also be chosen. In the simplest possible case, the network is modeled via a simple latency/bandwidth formula. For more complicated network models, many more than two parameters will be required. See 3.4 for a brief explanation of `SST/macro` network congestion models. A topology is also needed for constructing the network. In this case we choose a 2-D 4×4 torus (16 switches). The `topology.geometry` parameter takes an arbitrarily long list of numbers as the dimensions to the torus.

Finally, we must construct a node model. In this case, again, we use the simplest possible models for the node, network interface controller (NIC), and memory.

Parameter files can be constructed in a more modular way through the `include` statement. An alternative parameter file would be:

```
include machine.ini
# Launch parameters
node {
  app1 {
    indexing = block
    allocation = first_available
    launch_cmd = aprun -n2 -N1
    name = user_mpiapp_cxx
    argv =
    # Application parameters
    sendrecv_message_size = 128
  }
}
```

where in the first line we include the file `machine.ini`. All network, topology, and node parameters would be placed into a `machine.ini` file. In this way, multiple experiments can be linked to a common machine. Alternatively, multiple machines could be linked to the same application by creating and including an `application.ini`.

Using the deprecated (non-namespace) parameters the file would be:

```
# Launch parameters
launch_indexing = block
launch_allocation = first_available
launch_app1_cmd = aprun -n2 -N1
launch_app1 = user_mpiapp_cxx
launch_app1_argv =
# Network parameters
network_name = simple
network_bandwidth = 1.0GB/s
network_hop_latency = 100ns
```

```
# Topology – Ring of 4 nodes
topology_name = torus
topology_geometry = 4,4
# Node parameters
node_cores = 1
node_name = null
node_memory_model = null
nic_name = null
# Application parameters
sendrecv_message_size = 128
```

All of these are special keywords in the global namespace that get expanded into parameters in a specific namespace. We really, really do not recommend using these deprecated parameters anymore.

3.2 Abstract Machine Models

The preferred mode for usage of SST/macro will be through specifying parameters for well-defined abstract machine models. This represents an intermediate-level mode that should cover the vast majority of use cases. The highly configurable, detailed parameter files will remain valid but will represent advanced usage mode for corner cases. The primary advantage of the abstract machine models is a uniform set of parameters regardless of the underlying congestion model or accuracy level (Pisces or LogGOPSim). Each input file requires the usual set of software parameters given in 3.1. For hardware parameters, two initial parameters are required and one is optional.

```
congestion_model = pisces
amm_model = amm1
accuracy_parameter = 1024
```

Here we indicate the congestion model to be used (the packet-flow) and the overall machine model (abstract machine model #1). Currently valid values for the congestion model are **pisces** (most accurate, slowest) and **simple** (least accurate, fastest), but more congestion models should be supported in future versions. Currently valid values for the abstract machine model are **amm1**, **amm2**, **amm3**, see details below. Another model, **amm4**, that adds extra detail to the NIC and switches is pending and should be available soon. The details of individual abstract machine models are given in the following sections. The optional accuracy parameter is less well-defined and the exact meaning varies considerably between congestion models. In general, the accuracy parameter represents how coarse-grained the simulation is in bytes. It basically corresponds to a packet-size. How many bytes are modeled moving through the machine separately at a time? If the parameter is set to 8 bytes, e.g., that basically means we are doing flit-level modeling. If the parameter is set to 8192 bytes, e.g. that means we are doing very coarse-grained modeling which only really affects large messages. If the parameter is set to 100-1000 bytes, e.g., that means we are doing more fine-grained modeling on real packet sizes, but we are ignoring flit-level details.

3.2.1 Common Parameters

The following parameters define the CPU and compute power of the node (independent of memory subsystem). They are universal and are valid for all abstract machine models.

```
node {
  model = simple
  nsockets = 4
  proc {
    frequency = 2.1ghz
    ncores = 24
  }
}
```


3.2.2 AMM1

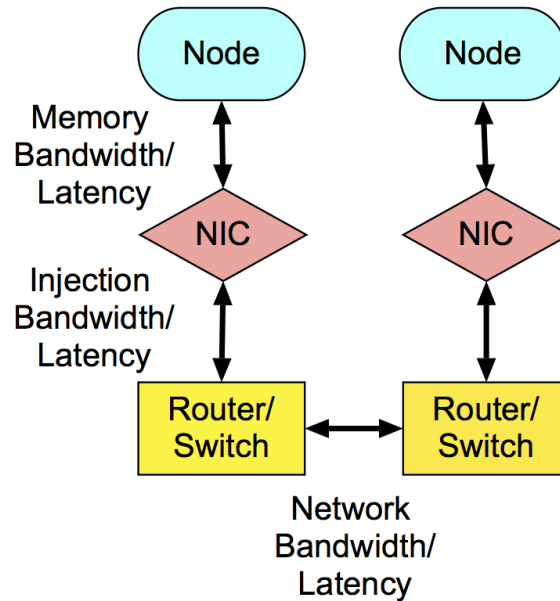


Figure 3.1: AMM1: Used when focusing on network traffic only

This is simplest abstract machine model and incorporates three basic components (i.e. congestion points). Each node has a memory subsystem and NIC (injection/ejection). Once packets are injected, they traverse a series of network switches. The memory, injection, and network are all defined by a bandwidth/latency parameter pair.

```
switch {  
  link {  
    bandwidth = 6GB/s  
    latency = 100ns  
  }  
}  
  
node {  
  nic {  
    injection {  
      bandwidth = 10GB/s  
      latency = 1us  
    }  
  }  
  memory {  
    bandwidth = 10GB/s  
    latency = 15ns  
  }  
}
```

The link latency is the latency required for a single packet to traverse one switch and hop to the next one in the network. Thus, even in the most basic of network models, there is still a notion of topology that affects the number of hops and therefore the latency. To compute the total network network latency as one would observe in an MPI ping-ping benchmark,

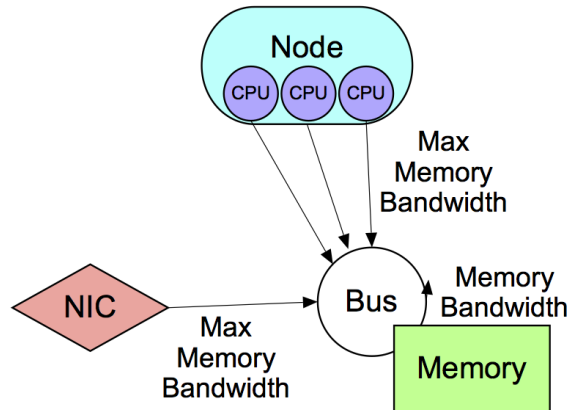


Figure 3.2: AMM2: Adds extra memory model details to AMM1

one would compute

$$lat = n_{hops} * lat_{hop} + 2 * lat_{inj}$$

using the hop latency and the injection latency.

This abstract machine model is a good place to start for getting a “lay of the land” for simulations - and the simplest to configure. However, it has a few deficiencies that can cause problems when there is serious memory or network congestion. More details (and their fixes) are given in the next abstract machine models.

3.2.3 AMM2

A major deficiency of AMM1 is that it grants exclusive access to memory resources. Two CPUs or the NIC cannot be using the memory subsystem in parallel. This is particularly problematic for large memory transfers (1 MB or greater). The memory system might be blocked for approx 1 ms, creating unphysical delays while other resources wait for access. A more realistic model allows multiple resources to access the memory, albeit with reduced bandwidth when congestion is observed. In many cases, multiple memory links or management units are connect to a shared bus. The bus determines to the total, aggregate memory bandwidth. However, the individual links determine the maximum observed bandwidth by any single component. AMM2 has all the same parameters as AMM1, but now allows an additional parameter for memory. These are special parameters used by the AMM configurations. They can be by-passed by directly using fully namespaced parameters (not shown).

```
node {
  memory {
    max_single_bandwidth = 5GB/s
    bandwidth = 10GB/s
    latency = 15ns
  }
}
```

The new parameter `max_single_bandwidth` now defines the maximum bandwidth any single component is allowed. Thus, even if the CPU is doing something memory intensive, 5 GB/s is still available to the NIC for network transfers. We remark

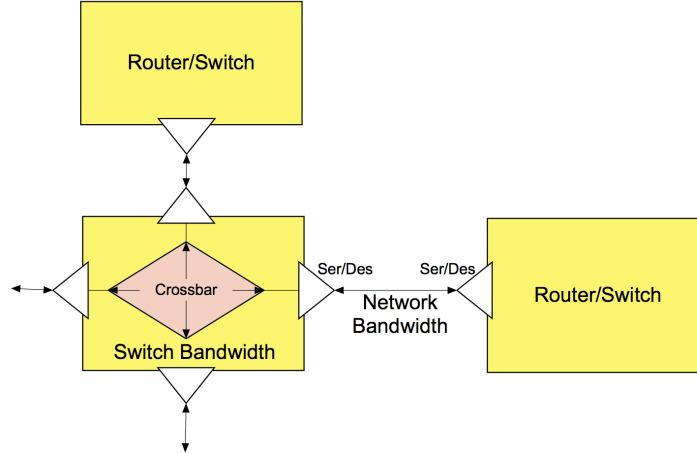


Figure 3.3: AMM3: Adds extra router (switch) details to AMM2

here that the memory parameters might be named something slightly more descriptive. However, as a rule, we want the AMM1 parameters to be a proper subset of the AMM2 parameters. Thus parameter names should not change - only new parameters should be added.

3.2.4 AMM3

A major deficiency of AMM2 is its inability to distinguish between the network link bandwidth (associated with the output port serializer/deserializer) and the switch bandwidth (associated with the crossbar that arbitrates packets). Only packets traveling the same path cause congestion on the network links in AMM1 and AMM2. However, packets “intersecting” at a switch - even if following separate paths - can cause congestion through sharing the switching fabric. AMM3 generalize the network parameters by adding a switch bandwidth. We note again here that AMM3 has all the same parameters as AMM2, plus the additional switch bandwidth parameter. Thus, higher-numbered abstract machine models always add more detail. These are special parameters used by the AMM configurations. They can be by-passed by directly using fully namespaced parameters (not shown) for more detailed configurations.

```
switch {
  xbar {
    bandwidth = 12GB/s
  }
  link {
    bandwidth = 6GB/s
    latency = 100ns
  }
}
```

3.3 Network Topologies and Routing

We here give a brief introduction to specifying different topologies and routing strategies. We will only discuss one basic example (torus). A more thorough introduction covering all topologies is planned for future releases. Excellent resources are “Principles and Practices of Interconnection Networks” by Brian Towles and William Dally published by Morgan Kaufman and “High Performance Datacenter Networks” by Dennis Abts and John Kim published by Morgan and Claypool.

3.3.1 Topology

Topologies are determined by two mandatory parameters.

```
topology.name = torus
topology.geometry = 4 4
```

Here we choose a 2D-torus topology with extent 4 in both the X and Y dimensions for a total of 16 nodes (Figure 3.4). The topology is laid out in a regular grid with network links connecting nearest neighbors. Additionally, wrap-around links connect the nodes on each boundary.

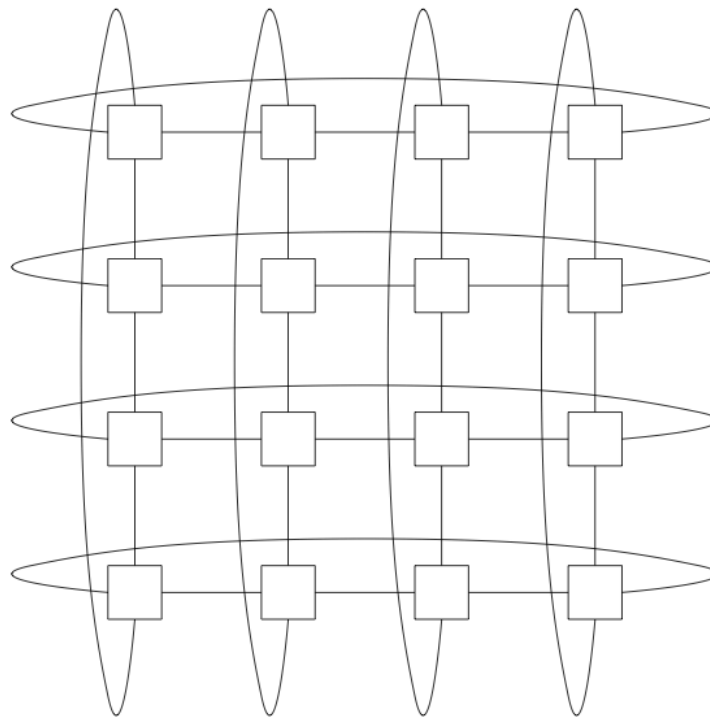


Figure 3.4: 4 x 4 2D Torus

The figure is actually an oversimplification. The `topology.geometry` parameter actually specifies the topology of the *network switches*, not the compute nodes. A torus is an example of a direct network in which each switch has one or more nodes

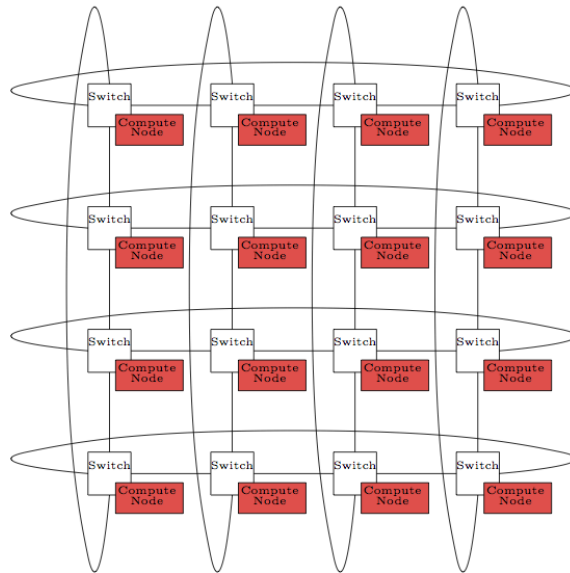


Figure 3.5: 4 x 4 2D Torus of Network Switches with Compute Nodes

“directly” connected to it. A more accurate picture of the network is given in Figure 3.5. While in many previous architectures there was generally a one-to-one correspondence between compute nodes and switches, more recent architectures have multiple compute nodes per switch (e.g. Cray Gemini with two nodes, Cray Aries with four nodes). Multiple nodes per switch can be specified via a concentration parameter:

```
topology {
  name = torus
  geometry = 4 4
  concentration = 2
}
```

which would now generate a torus topology with 16 switches and 32 compute nodes.

Another subtle modification of torus (and other networks) can be controlled by giving the X , Y , and Z directions different bandwidth. The above network could be modified as

```
topology {
  name = torus
  geometry = 4 4
  redundant = 2 1
}
```

giving the the X -dimension twice the bandwidth of the Y -dimension. This pattern DOES exist in some interconnects as a load-balancing strategy. A very subtle point arises here. Consider two different networks:

```
topology {
  name = torus
  geometry = 4 4
  redundant = 1 1
}
switch.link.bandwidth = 2GB/s
```

```

topology {
  name = torus
  geometry = 4 4
  redundant = 2 2
}
switch.link.bandwidth = 1GB/s

```

For some coarse-grained models, these two networks are exactly equivalent. In more fine-grained models, however, these are actually two different networks. The first network has ONE link carrying 2 GB/s. The second network has TWO links each carrying 1 GB/s.

3.3.2 Routing

By default, SST/macro uses the simplest possible routing algorithm: dimension-order minimal routing (Figure 3.6). In going

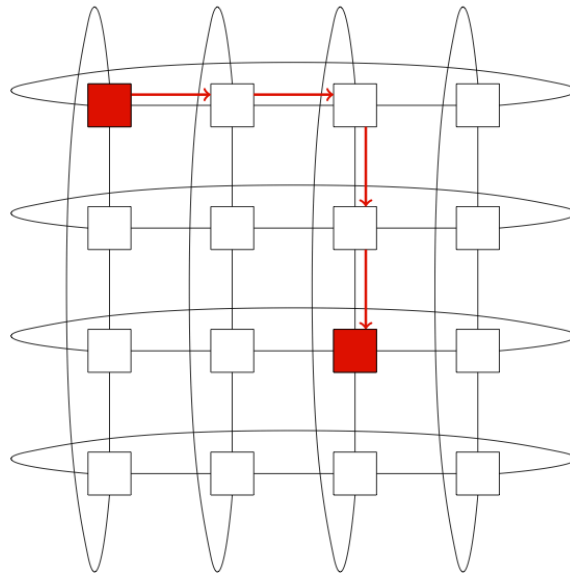


Figure 3.6: Dimension-Order Minimal Routing on a 2D Torus

from source to destination, the message first travels along the X -dimension and then travels along the Y -dimension. The above scheme is entirely static, making no adjustments to avoid congestion in the network. SST/macro supports a variety of adaptive routing algorithms. This can be specified:

```

switch {
  router {
    name = min_ad
  }
}

```

which specifies minimal adaptive routing. There are now multiple valid paths between network endpoints, one of which is illustrated in Figure 3.7. At each network hop, the router chooses the *productive* path with least congestion. In some cases,




Figure 3.7: Adaptive Minimal Routing on a 2D Torus

however, there is only one minimal path (node (0,0) sending to (2,0) with only X different). For these messages, minimal adaptive is exactly equivalent to dimension-order routing. Other supported routing schemes are valiant and UGAL. More routing schemes are scheduled to be added in future versions. A full description of more complicated routing schemes will be given in its own chapter in future versions. For now, we direct users to existing resources such as “High Performance Datacenter Networks” by Dennis Abts and John Kim.

3.4 Network Model

Network models can be divided into several categories. SST/macro supports analytic models, which estimate network delays via basic latency/bandwidth formulas, and packet models, which model step-by-step the transit of individuals through the interconnect. A third class of models (flow models), was previously supported but are now discontinued due to the much better scalability of packet models.

3.4.1 Analytic Models: MACRELS

The analytic models in SST/macro are colloquially referred to as MACRELS (MTL for AnalytiC REally Lightweight Simulation). The MTL (message transfer layer) moves entire network flows from point-to-point without packetizing them into smaller chunks. Thus an entire 1 MB MPI message is transported as a single chunk of data. The majority of MACRELS models are based on the LogP set of approximations:

$$\Delta t = \alpha + \beta N$$

where Δt is the time delay, α is the minimum latency of the communication, β is the inverse bandwidth (s/B), and N is the number of bytes. In abstract machine models, these methods are selected as:

```
congestion_model = logP
```

3.4.2 Packet Models: PISCES

PISCES (Packet-flow Interconnect Simulation for Congestion at Extreme Scale) breaks network flows (MPI messages) into individual packets and models each packet individually. In abstract machine models, PISCES can be selected as:

```
congestion_model = pscs
```

In reality, packets are further subdivided into flits (flow-control units). Flit-level detail would be way too computationally intense for large-scale simulation. All routing decisions are made on packets as a whole. Two flits in the same packet cannot take different paths through the network. However, they may not travel together. PISCES provides two mechanisms for treating flit-level flow control.

PISCES simple model

In the simple model, each router uses a basic store-and-forward mechanism. Flits are not allowed to “separate” and always travel as a single unit. The entire packet has to be stored within a router before it can be forwarded to the next router. The simple model affects the arbitrator that decided when and how to transmit flits. To select a simple model:

```
arbitrator = simple
```

The simple model is the least computationally expensive. However, for large packet sizes, it can produce erroneously high latencies. To tune the packet size for abstract machine models, set:

```
accuracy_parameter = 1024B
```

or equivalently

```
mtu = 1024B
```

which sets the packet size to 1024B. For the simple model, packet sizes larger than 256-512B are not recommended. Packet sizes on production supercomputers are often small (96-128B). Small packet sizes with the simple model can be a good compromise for having more fine-grained routing but cheaper congestion modeling in the arbitrator.

PISCES cut-through model

In the cut-through model, routing decisions still occur at the packet-level. However, some attempt is made to account for pipelining of flits across different router stages. Somewhat similar to the LogP models used above, latency/bandwidth formulas are used to estimate packet delays. However, the cut-through model adds more details. It’s requested as:

```
arbitrator = cut_through
```

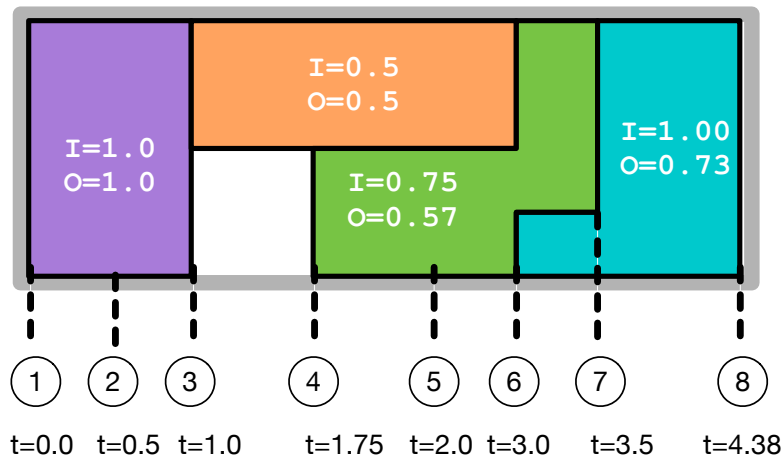



Figure 3.8: Timeline of four different packets passing through a PISCES cut-through bandwidth arbitrator. The incoming bandwidth (I) and outgoing bandwidth (O) are shown for each packet. Time is the horizontal axis. Bandwidth consumed by a packet is shown by the vertical extent of each packet. The individual events are 1) First packet arrives 2) Second packet arrives with reduced bandwidth but no available bandwidth 3) First packet finishes. Second packet can begin sending. 4) Third packet arrives and begins sending with remaining bandwidth. 5) Fourth packet arrives, but no available bandwidth. 6) Second packet finishes. Third packet increases bandwidth. Fourth packet can begin sending. 7) Third packet finishes. Fourth packet increases bandwidth. 8) Fourth packet finishes. Full details are given in the text.

Figure 3.8 shows a timeline for the data being transmitted through a crossbar, SerDes, or other network component with a “fixed bandwidth.” Each component is essentially a pipe with some flow bandwidth. The arbitrator divides its limited bandwidth amongst incoming packets. Packets fill the pipeline, consuming bandwidth. In contrast to the completely discrete simple model, packets can “multiplex” in the component sharing an arbitrary bandwidth partition. Modeling a packet delay starts with two input parameters and computes three output parameters.

- A : Packet head arrival (input)
- I : Packet incoming bandwidth (input)
- H : Packet head departure (output)
- T : Packet tail departure (output)
- O : Packet outgoing bandwidth (output)

In the simple model, a packet either consumes all the bandwidth or none of the bandwidth. To account for flit-level pipelining, the cut-through model allows packets to consume partial bandwidths. Consider an arbitrator that has a maximum bandwidth of 1.0. The first packet (purple, Figure 3.8) arrives with a full incoming bandwidth of 1.0 and head arrival of $t=0.0$. It therefore consumes all the available bandwidth. The head of the packet can actually leave immediately (as it must to properly pipeline or cut-through). The tail leaves after all bytes have sent at $t=1.0$. Thus for the first packet we have $H=0.0$, $T=1.0$, and $O=1.0$.

The second packet (orange) arrives at $t=0.5$. Upon arrival there is no bandwidth available as the first packet is consuming the maximum. Only after the first packet finishes can the second packet begin. The second packet arrives and leaves with a reduced bandwidth of 0.5. Thus we have $H=1.0$, $T=3.0$, and $O=0.5$.

The third packet (green) arrives at $t=1.75$. Upon arrival there is some bandwidth, but not enough to match the incoming bandwidth of 0.57. Thus the third packet is slowed initially. After the second packet finished, the third packet can send at increased bandwidth. The computation here is a bit more complicated. Packet 3 can actually consume MORE than 0.6 bandwidth units. Between steps 4 and 6, packet 3 has accumulated in a local buffer in the router. Thus even though the incoming bandwidth is only 0.6, there are several flits that are available to send immediately at full bandwidth waiting in the buffer. Thus results in an effective bandwidth of 0.75 for the remainder of the packet’s time in the arbitrator. Thus we end up with $H=1.75$, $T=3.5$, and $O=0.57$. Even though the packet is initially delayed, the buffers compensate for the delay and allow the outgoing bandwidth to “catch up” with the incoming bandwidth.

Finally, the fourth packet (blue) arrives at $t=3.0$. There is some available bandwidth. After the third packet finishes, the fourth packet can now send at maximum. Because of the initial delay, the outgoing bandwidth is somewhat reduced. We have $H=3.0$, $T=4.38$, and $O=0.73$.

3.4.3 Flow

The flow model, in simple cases, corrects the most severe problems of the packet model. Instead of discrete chunks, messages are modeled as fluid flows moving through the network. Congestion is treated as a fluid dynamics problem, sharing bandwidth between competing flows. In contrast to LogP models, flow models can account fairly well for congestion. Without congestion, a flow only requires a FLOW START and FLOW STOP event to be modeled (see tutorial on discrete event simulation in 3.7). While the packet model would require many, many events to simulate a 1 MB message, the flow model might only require two. With congestion, flow update events must be scheduled whenever congestion changes on a network link. For

limited congestion, only a few update events must occur. The flow model also corrects the latency and multiplexing problems in the PISCES simple model, providing higher-accuracy for coarse-grained simulation.

The flow model starts to break down for large systems or under heavy congestion. In the packet model, all congestion events are “local” to a given router. The number of events is also constant in packet models regardless of congestion since we are modeling a fixed number of discrete units. In flow models, flow update events can be “non-local,” propagating across the system and causing flow update events on other routers. When congestion occurs, this “ripple effect” can cause the number of events to explode, overwhelming the simulator. For large systems or heavy congestion, the flow model is actually much slower than the packet model. Support for this model has been completely removed.

3.5 Basic MPI Program

Let us go back to the simple send/rcv skeleton and actually look at the code. This code should be compiled with SST compiler wrappers installed in the `bin` folder.

```
1  #include <stdlib.h>
2  #include <stdio.h>
3  #include <mpi.h>
4
5  #define sstmac_app_name simple_test
6
7  int main(int argc, char **argv)
8  {
9      int message_size = 128;
10     int me, nproc;
11     int tag = 0;
12     int dst = 1;
13     int src = 0;
14     MPI_Status stat;
15
16     MPI_Init(&argc,&argv);
17     MPI_Comm world = MPI_COMM_WORLD;
18     MPI_Comm_rank(world,&me);
19     MPI_Comm_size(world,&nproc);
```

The starting point is creating a main routine for the application. The simulator itself already provides a `main` routine. The SST compiler automatically changes the function name to `user_skeleton_main`, which provides an entry point for the application to actually begin. When SST/macro launches, it will invoke this routine and pass in any command line arguments specified via the `app1.argv` parameter. Upon entering the main routine, the code is now indistinguishable from regular MPI C++ code. In the parameter file to be used with the simulation, you must set

```
node.app1.name = simple_test
```

The name associated to the application is given by the `sstmac_app_name` macro. This macro must be defined to a unique string name in the source file containing `main`. SST/macro will automatically associate the given main routine with the string internally. That application can then be selected in the input file with `app1.name`.

At the very top of the file, the `mpi.h` header is actually mapped by the SST compiler to an SST/macro header file. This header provides the MPI API and configures MPI function calls to link to SST/macro instead of the real MPI library. The code now proceeds:

```
1  if (nproc != 2) {
2      fprintf(stderr, "sendrecv only runs with two processors\n");
3      abort();
4  }
```

```

5  if (me == 0) {
6      MPI_Send(NULL, message_size, MPI_INT, dst, tag, world);
7      printf("rank %i sending a message\n", me);
8  }
9  else {
10     MPI_Recv(NULL, message_size, MPI_INT, src, tag, world, &stat);
11     printf("rank %i receiving a message\n", me);
12 }
13 MPI_Finalize();
14 return 0;
15 }

```

Here the code just checks the MPI rank and sends (rank 0) or receives (rank 1) a message.

3.6 Launching, Allocation, and Indexing

3.6.1 Launch Commands

Just as jobs must be launched on a shared supercomputer using Slurm or aprun, SST/macro requires the user to specify a launch command for the application. Currently, we encourage the user to use aprun from Cray, for which documentation can easily be found online. In the parameter file you specify, e.g.

```

node {
  app1 {
    name = user_mpiapp_cxx
    launch_cmd = aprun -n 8 -N 2
  }
}

```

which launches an external user C++ application with eight ranks and two ranks per node. The aprun command has many command line options (see online documentation), some of which may be supported in future versions of SST/macro. In particular, we are in the process of adding support for thread affinity, OpenMP thread allocation, and NUMA containment flags. Most flags, if included, will simply be ignored.

3.6.2 Allocation Schemes

In order for a job to launch, it must first allocate nodes to run on. Here we choose a simple 2D torus

```

topology.name = torus
topology.geometry = 3 3
topology.concentration = 1

```

which has 9 nodes arranged in a 3x3 mesh. For the launch command `aprun -n 8 -N 2`, we must allocate 4 compute nodes from the pool of 9. Our first option is to specify the first available allocation scheme (Figure 3.9)

```

node.app1.allocation = first_available

```

In first available, the allocator simply loops through the list of available nodes as they are numbered by the topology object. In the case of a 2D torus, the topology numbers by looping through columns in a row. In general, first available will give a contiguous allocation, but it won't necessarily be ideally structured.

To give more structure to the allocation, a Cartesian allocator can be used (Figure 3.10).

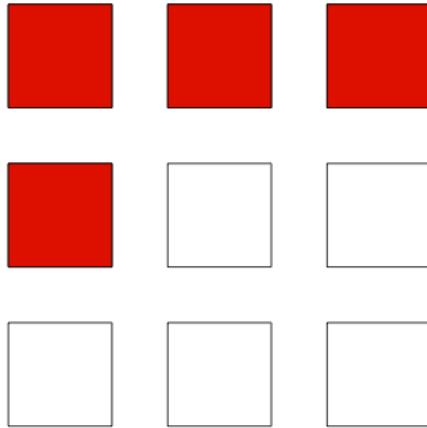


Figure 3.9: First available Allocation of 4 Compute Codes on a 3x3 2D Torus

```

appl {
  allocation = cartesian
  cart_sizes = 2 2
  cart_offsets = 0 0
}

```

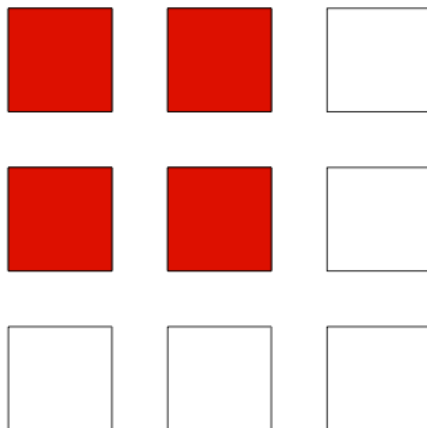


Figure 3.10: Cartesian Allocation of 4 Compute Codes on a 3x3 2D Torus

Rather than just looping through the list of available nodes, we explicitly allocate a 2x2 block from the torus. If testing how “topology agnostic” your application is, you can also choose a random allocation.

```

node.appl.allocation = random

```

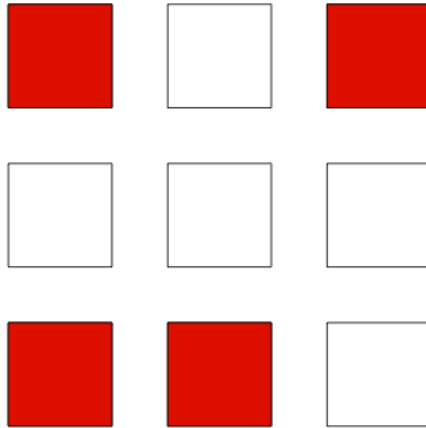


Figure 3.11: Random Allocation of 4 Compute Codes on a 3x3 2D Torus

In many use cases, the number of allocated nodes equals the total number of nodes in the machine. In this case, all allocation strategies allocate the same *set* of nodes, i.e. the whole machine. However, results may still differ slightly since the allocation strategies still assign an initial numbering of the node, which means a random allocation will give different results from Cartesian and first available.

Indexing Schemes

Once nodes are allocated, the MPI ranks (or equivalent) must be assigned to physical nodes, i.e. indexed. The simplest strategies are block and round-robin. If only running one MPI rank per node, the two strategies are equivalent, indexing MPI ranks in the order received from the allocation list. If running multiple MPI ranks per node, block indexing tries to keep consecutive MPI ranks on the same node (Figure 3.12).

```
node.app1.indexing = block
```

In contrast, round-robin spreads out MPI ranks by assigning consecutive MPI ranks on different nodes (Figure 3.13).

```
node.app1.indexing = round_robin
```

Finally, one may also choose

```
node.app1.indexing = random
```

Random allocation with random indexing is somewhat redundant. Random allocation with block indexing is *not* similar to Cartesian allocation with random indexing. Random indexing on a Cartesian allocation still gives a contiguous block of nodes, even if consecutive MPI ranks are scattered around. A random allocation (unless allocating the whole machine) will not give a contiguous set of nodes.

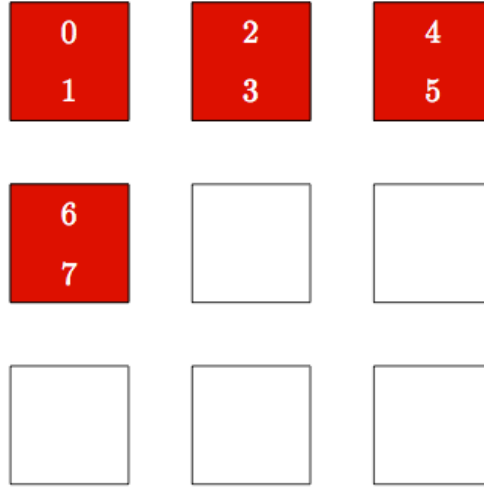


Figure 3.12: Block Indexing of 8 MPI Ranks on 4 Compute Nodes

3.7 Discrete Event Simulation

Although not necessary for using the simulator, a basic understanding of discrete event simulation can be helpful in giving users an intuition for network models and parameters. Here we walk through a basic program that executes a single send/recv pair. SST/macro simulates many parallel processes, but itself runs as a single process with only one address space (SST/macro can actually run in parallel mode, but we ignore that complication here). SST/macro manages each parallel process as a user-space thread (application thread), allocating a thread stack and frame of execution. User-space threading is necessary for large simulations since otherwise the kernel would be overwhelmed scheduling thousands of threads.

SST/macro is driven by a simulation thread which manages the user-space thread scheduling (Figure 3.14). In the most common (and simplest) use case, all user-space threads are serialized, running one at a time. The main simulation thread must manage all synchronizations, yielding execution to process threads at the appropriate times. The main simulation thread is usually abbreviated as the DES (discrete event simulation) thread. The simulation progresses by scheduling future events. For example, if a message is estimated to take $5\ \mu s$ to arrive, the simulator will schedule a MESSAGE ARRIVED event $5\ \mu s$ ahead of the current time stamp. Every simulation starts by scheduling the same set of events: launch process 0, launch process 1, etc.

The simulation begins at time $t = 0\ \mu s$. The simulation thread runs the first event, launching process 0. The context of process 0 is switched in, and SST/macro proceeds running code as if it were actually process 0. Process 0 starts a blocking send in Event 1. For process 0 to perform a send in the simulator, it must *schedule* the necessary events to simulate the send. Most users of SST/macro will never need to explicitly schedule events. Discrete event details are always hidden by the API and executed inside library functions. In this simple case, the simulator estimates the blocking send will take $1\ \mu s$. It therefore schedules a SEND DONE (Event 4) $1\ \mu s$ into the future before blocking. When process 0 blocks, it yields execution back to the main simulation.

At this point, no time has yet progressed in the simulator. The DES thread runs the next event, launching process 1, which executes a blocking receive (Event 3). Unlike the blocking send case, the blocking receive does not schedule any events. It

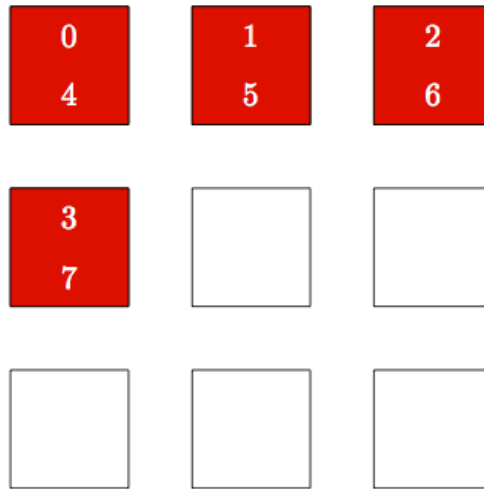


Figure 3.13: Round-Robin Indexing of 8 MPI Ranks on 4 Compute Nodes

	Sim Thread	Process 0	Process 1
$t = 0\mu s$	0)Launch proc 0 2)Launch proc 1	1)Block until send complete	3)Post recv to NIC; block
$t = 1\mu s$	4)Send done; unblock proc 0 6)Deliver msg to NIC 1 ($1\mu s$)	5)Wait for ack; block	
$t = 2\mu s$	7)Recv at NIC 1; unblock proc 1		8)Send ack for recv ($1\mu s$); block
$t = 3\mu s$	9)Deliver ack to NIC 0 ($1\mu s$) 10)Send done; unblock proc 1		11)Continue execution...
$t = 4\mu s$	12)Recv at NIC 0; unblock proc 0	13)Continue execution...	

Figure 3.14: Progression of Discrete Event Simulation for Simple Send/Recv Example

cannot know when the message will arrive and therefore blocks without scheduling a RECV DONE event. Process 1 just registers the receive and yields back to the DES thread.

At this point, the simulator has no events left at $t=0\mu s$ and so it must progress its time stamp. The next event (Event 4)

is SEND DONE at $t=1\ \mu\text{s}$. The event does two things. First, now that the message has been injected into the network, the simulator estimates when it will arrive at the NIC of process 1. In this case, it estimates $1\ \mu\text{s}$ and therefore schedules a MESSAGE ARRIVED event in the future at $t=2\ \mu\text{s}$ (Event 7). Second, the DES thread unblocks process 0, resuming execution of its thread context. Process 0 now posts a blocking receive, waiting for process 1 to acknowledge receipt of its message.

The simulator is now out of events at $t=1\ \mu\text{s}$ and therefore progresses its time stamp to $t=2\ \mu\text{s}$. The message arrives (Event 7), allowing process 1 to complete its receive and unblock. The DES thread yields execution back to process 1, which now executes a blocking send to ack receipt of the message. It therefore schedules a SEND DONE event $1\ \mu\text{s}$ in the future (Event 10) and blocks, yielding back to the DES thread. This flow of events continues until all the application threads have terminated. The DES thread will run out of events, bringing the simulation to an end.

3.8 Using DUMPI

3.8.1 Building DUMPI

As noted in the introduction, SST/macro is primarily intended to be an on-line simulator. Real application code runs, but SST/macro intercepts calls to communication (MPI) and computation functions to simulate time passing. However, SST/macro can also run off-line, replaying application traces collected from real production runs. This trace collection and trace replay library is called DUMPI.

Although DUMPI is automatically included as a subproject in the SST/macro download, trace collection can be easier if DUMPI is built independently from SST/macro. The code can be downloaded from <https://bitbucket.org/sst-ca/dumpi>. If downloaded through Mercurial, one must initialize the build system and create the configure script.

```
dumpi> ./bootstrap.sh
```

DUMPI must be built with an MPI compiler.

```
dumpi/build> ../configure CC=mpicc CXX=mpicxx \
--enable-libdumpi --prefix=$DUMPI_PATH
```

The `--enable-libdumpi` flag is needed to configure the trace collection library. After compiling and installing, a `libdumpi` will be added to `$DUMPI_PATH/lib`.

Collecting application traces requires only a trivial modification to the standard MPI build. Using the same compiler, simply add the DUMPI library path and library name to your project's `LDFLAGS`.

```
your_project/build> ../configure CC=mpicc CXX=mpicxx \
LDFLAGS='-L$DUMPI_PATH/lib -ldumpi'
```

3.8.2 Trace Collection

DUMPI works by overriding *weak symbols* in the MPI library. In all MPI libraries, functions such as `MPI_Send` are only weak symbol wrappers to the actual function `PMPI_Send`. DUMPI overrides the weak symbols by implementing functions with the symbol `MPI_Send`. If a linker encounters a weak symbol and regular symbol with the same name, it ignores the weak symbol. DUMPI functions look like

```

1 int MPI_Send(...)
2 {
3     /** Start profiling work */
4     ...
5     int rc = PMPI_Send(...);
6     /** Finish profiling work */
7     ...
8     return rc;
9 }

```

collecting profile information and then directly calling the PMPI functions.

We examine DUMPI using a very basic example program.

```

1 #include <mpi.h>
2 int main(int argc, char** argv)
3 {
4     MPI_Init(&argc, &argv);
5     MPI_Finalize();
6     return 0;
7 }

```

After compiling the program named `test` with DUMPI, we run MPI in the standard way.

```
example> mpiexec -n 2 ./test
```

After running, there are now three new files in the directory.

```

example> ls dumpi*
dumpi-2013.09.26.10.55.53-0000.bin
dumpi-2013.09.26.10.55.53-0001.bin
dumpi-2013.09.26.10.55.53.meta

```

DUMPI automatically assigns a unique name to the files from a timestamp. The first two files are the DUMPI binary files storing separate traces for MPI rank 0 and rank 1. The contents of the binary files can be displayed in human-readable form by running the `dumpi2ascii` program, which should have been installed in `$DUMPI_PATH/bin`.

```
example> dumpi2ascii dumpi-2013.09.26.10.55.53-0000.bin
```

This produces the output

```

MPI_Init entering at walltime 8153.0493, cputime 0.0044 seconds in thread 0.
MPI_Init returning at walltime 8153.0493, cputime 0.0044 seconds in thread 0.
MPI_Finalize entering at walltime 8153.0493, cputime 0.0045 seconds in thread 0.
MPI_Finalize returning at walltime 8153.0498, cputime 0.0049 seconds in thread 0.

```

The third file is just a small metadata file DUMPI used to configure trace replay.

```

hostname=deepthought.magrathea.gov
numprocs=2
username=slartibartfast
starttime=1380218153
fileprefix=dumpi-2013.09.26.10.55.53
version=1
subversion=1
subsubversion=0

```

3.8.3 Trace Replay

It is often useful to validate the correctness of a trace. Sometimes there can be problems with trace collection. There are also a few nooks and crannies of the MPI standard left unimplemented. To validate the trace, you can run in a special debug mode that runs the simulation with a very coarse-grained model to ensure as quickly as possible that all functions execute correctly. This can be done straightforwardly by running the executable with the dumpi flag: `sstmac --dumpi`.

To replay a trace in the simulator, a small modification is required to the example input file in 3.1. We have two choices for the trace replay. First, we can attempt to *exactly* replay the trace as it ran on the host machine. Second, we could replay the trace on a new machine or different layout.

For exact replay, the key issue is specifying the machine topology. For some architectures, topology information can be directly encoded into the trace. This is generally true on Blue Gene, but not Cray. When topology information is recorded, trace replay is much easier. The parameter file then becomes, e.g.

```
node {
  app1 {
    launch_type = dumpi
    indexing = dumpi
    allocation = dumpi
    name = parsedumpi
    dumpi_metaname = testbgp.meta
  }
}
```

We have a new parameter `launch_type` set to `dumpi`. This was implicit before, taking the default value of `skeleton`. We also set indexing and allocation parameters to read from the DUMPI trace. The application name is a special app that parses the DUMPI trace. Finally, we direct SST/macro to the DUMPI metafile produced when the trace was collected. To extract the topology information, locate the `.bin` file corresponding to MPI rank 0. To print topology info, run

```
traces> dumpi2ascii -H testbgp-0000.bin
```

which produces the output

```
version=1.1.0
starttime=Fri Nov 22 13:53:58 2013
hostname=R00-M1-N01-J01.challenger
username=<none>
meshdim=3
meshsize=[4, 2, 2]
meshcrd=[0, 0, 0]
```

Here we see that the topology is 3D with extent 4,2,2 in the X,Y,Z directions. At present, the user must still specify the topology in the parameter file. Even though SST/macro can read the topology *dimensions* from the trace file, it cannot read the topology *type*. It could be a torus, dragonfly, or fat tree. The parameter file therefore needs

```
topology {
  name = torus
  geometry = 4 2 2
}
```

Beyond the topology, the user must also specify the machine model with bandwidth and latency parameters. Again, this is information that cannot be automatically encoded in the trace. It must be determined via small benchmarks like ping-pong. An example file can be found in the test suite in `tests/test_configs/testdumpibgp.ini`.

If no topology info could be recorded in the trace, more work is needed. The only information recorded in the trace is the hostname of each MPI rank. The parameters are almost the same, but with allocation now set to `hostname`. Since no

topology info is contained in the trace, a hostname map must be put into a text file that maps a hostname to the topology coordinates. The new parameter file, for a fictional machine called deep thought

```
# Launch parameters
node {
  app1 {
    launch_type = dumpi
    indexing = dumpi
    allocation = hostname
    name = parsedumpi
    dumpi_metaname = dumpi-2013.09.26.10.55.53.meta
    dumpi_mapname = deepthought.map
  }
}
# Machine parameters
topology {
  name = torus
  geometry = 2 2
}
```

In this case, we assume a 2D torus with four nodes. Again, DUMPI records the hostname of each MPI rank during trace collection. In order to replay the trace, the mapping of hostname to coordinates must be given in a node map file, specified by the parameter `launch_dumpi_mapname`. The node map file has the format

```
4 2
nid0 0 0
nid1 0 1
nid2 1 0
nid3 1 1
```

where the first line gives the number of nodes and number of coordinates, respectively. Each hostname and its topology coordinates must then be specified. More details on building hostname maps are given below.

We can also use the trace to experiment with new topologies to see performance changes. Suppose we want to test a crossbar topology.

```
# Launch parameters
node {
  app1 {
    indexing = block
    allocation = first_available
    dumpi_metaname = dumpi-2013.09.26.10.55.53.meta
    name = parsedumpi
    size = 2
  }
}
# Machine parameters
topology {
  name = crossbar
  geometry = 4
}
```

We no longer use the DUMPI allocation and indexing. We also no longer require a hostname map. The trace is only used to generate MPI events and no topology or hostname data is used. The MPI ranks are mapped to physical nodes entirely independent of the trace.

3.9 Using Score-P and OTF2 (Beta)

OTF2 is part of Score-P. Sources for both can be found [here](#)

```
http://www.vi-hps.org/projects/score-p
```

Trace collection requires both Score-P and OTF2 installations. Trace replay with SST/macro requires OTF2.

3.9.1 Trace Collection

Score-P's default collection strategy will include every function call in the trace, making even small programs produce untenably large traces. Score-P supports collection filters, which can restrict collection at a minimum to MPI and OMP function calls. At the end of the program's runtime, traces from each rank are put in a common directory. An MPI program must be compiled with Score-P to produce traces:

```
scorep-mpicxx -o test.exe test.cc
```

To limit the size of the traces, run the program with:

```
# these environment variables are picked up by Score-P at runtime
export SCOREP_ENABLE_TRACING=true
export SCOREP_TOTAL_MEMORY=1G
export SCOREP_FILTERING_FILE='scorep.filter'

mpirun -n 2 test.exe
```

The file `scorep.filter` should contain:

```
SCOREP_REGION_NAMES_BEGIN EXCLUDE *
```

To view a plain-text representation of the trace after running, use the `otf2-print` tool.

```
otf2-print scorep-*/traces.otf2
```

3.9.2 Trace Replay

SST/macro will use a trace replay skeleton for OTF2 in much the same way as it does for dumpi. SST/macro trace replays configured using *.ini files.

```
...
node {
  app1 {
    otf2_timescale = 1.0
    name = otf2_trace_replay_app
    size = N
    otf2_metafile = <trace-root>/scorep-20170309_1421_27095992608015568/traces.otf2
    # debugging output
    otf2_print_mpi_calls=false
    otf2_print_trace_events=false
    otf2_print_time_deltas=false
    otf2_warn_unknown_callback=false
  }
}
```

3.10 Call Graph Visualization

Generating call graphs requires a special build of SST/macro.

```
build> ../configure --prefix=$INSTALL_PATH --enable-graphviz
```

The `--enable-graphviz` flag defines an instrumentation macro throughout the SST/macro code. This instrumentation must be *compiled* into SST/macro. In the default build, the instrumentation is not added since the instrumentation has a high overhead. However, SST/macro only instruments a select group of the most important functions so the overhead should only be 10-50%. After installing the instrumented version of SST/macro, a call graph is collected by adding a simple filename to the parameter file.

```
call_graph = <fileroot>
```

After running, a `<fileroot>.callgrind.out` file should appear in the folder.

To visualize the call graph, you must download KCachegrind: <http://kcachegrind.sourceforge.net/html/Download.html>. KCachegrind is built on the KDE environment, which is simple to build for Linux but can be very tedious for Mac. The download also includes a QKachegrind subfolder, providing the same functionality built on top of Qt. This is highly recommended for Mac users.

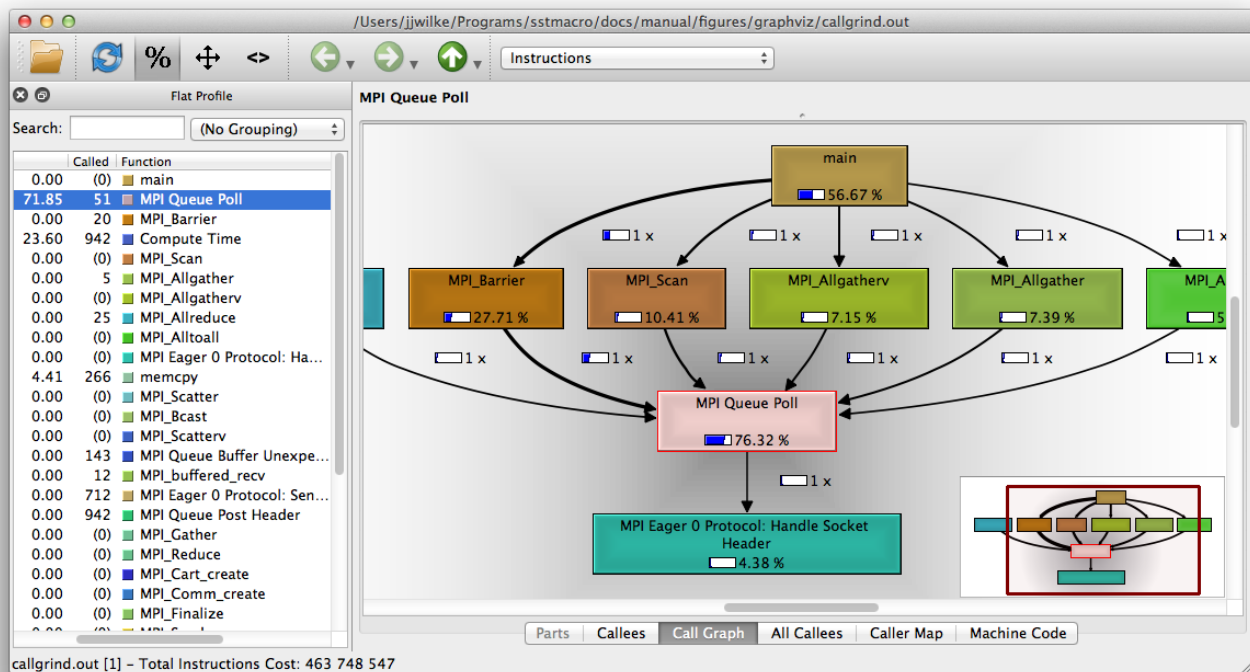


Figure 3.15: QKachegrind GUI

The basic QCachegrind GUI is shown in Figure 3.15. On the left, a sidebar contains the list of all functions instrumented with the percent of total execution time spent in the function. In the center pane, the call graph is shown. To navigate the call graph, a small window in the bottom right corner can be used to change the view pane. Zooming into one region (Figure 3.16), we see a set of MPI functions (Barrier, Scan, Allgather). Each of the functions enters a polling loop, which dominates the total execution time. A small portion of the polling loop calls the “Handle Socket Header” function. Double-clicking this node unrolls more details in the call graph (Figure 3.17). Here we see the function splits execution time between buffering messages (memcpy) and posting headers (Compute Time).

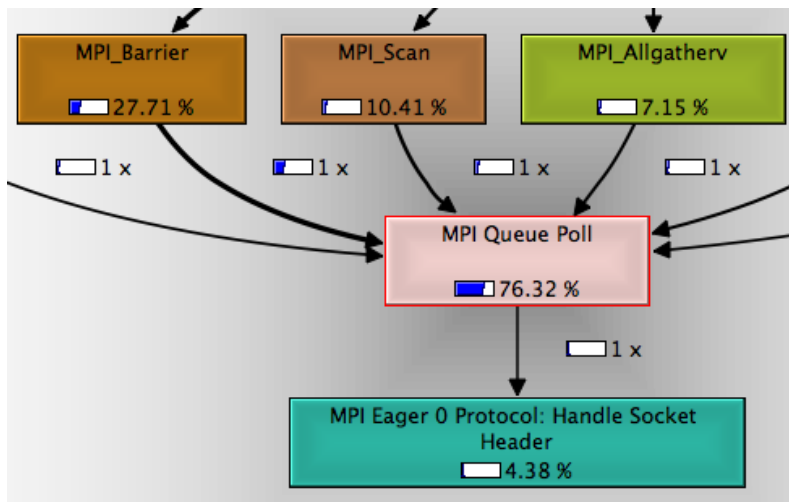


Figure 3.16: QCachegrind Call Graph of MPI Functions

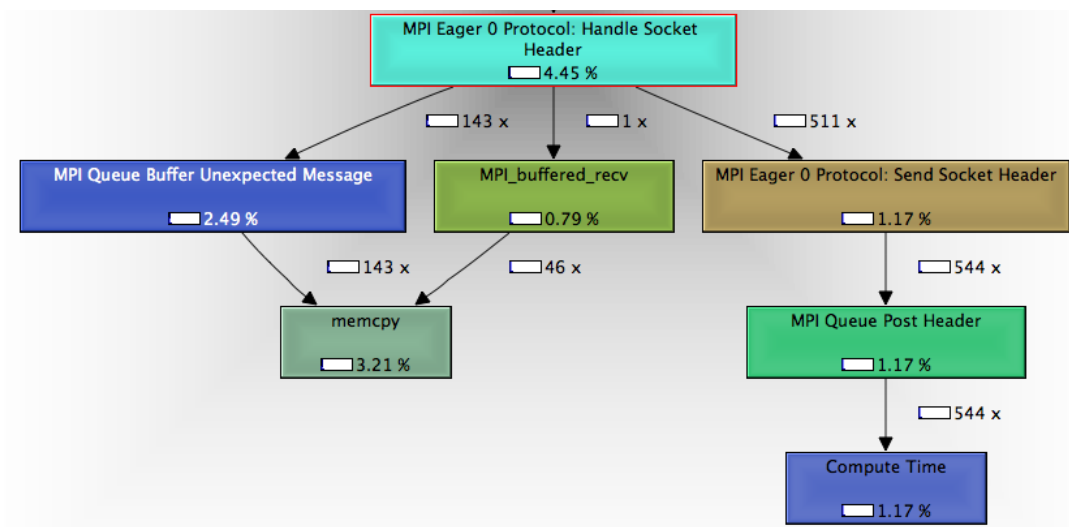


Figure 3.17: QCachegrind Expanded Call Graph of Eager 0 Function

3.11 Spyplot Diagrams

Spyplots visualize communication matrices, showing either the number of messages or number of bytes sent between two network endpoints. They are essentially contour diagrams, where instead of a continuous function $F(x, y)$ we are plotting the communication matrix $M(i, j)$. An example spyplot is shown for a simple application that only executes an `MPI_Allreduce` (Figure 3.18). Larger amounts of data (red) are sent to nearest neighbors while decreasing amounts (blue) are sent to MPI ranks further away.

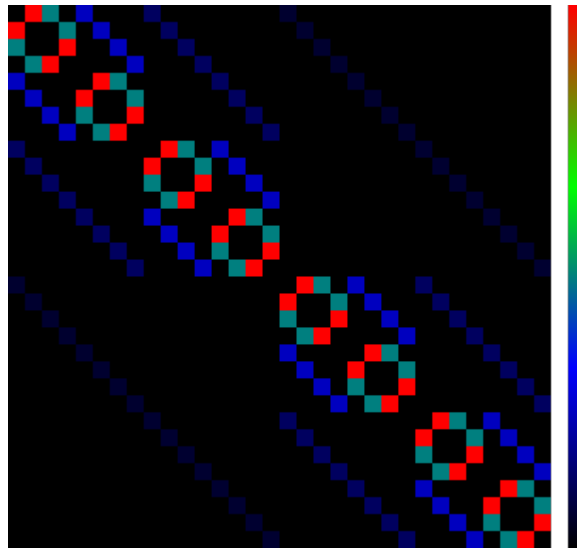


Figure 3.18: Spyplot of Bytes Transferred Between MPI Ranks for `MPI_Allreduce`

Various spyplots can be activated by boolean parameters in the input file. The most commonly used are the MPI spyplots, for which you must add

```
mpi_spyplot = <fileroot>
```

After running there will be a .csv and .png file in the folder, with e.g. `fileroot = test`

```
example> ls
test.png
test.csv
```

`test.png` shows the number of bytes exchanged between MPI ranks. To extend the analysis you can specify

```
network_spyplot = <fileroot>
```

A new csv/png will appear showing the number of bytes exchanged between physical nodes, accumulating together all MPI ranks sharing the same node. This gives a better sense of spatial locality when many MPI ranks are on the same node.




Figure 3.19: Application Activity (Fixed-Time Quanta; FTQ) for Simple MPI Test Suite

3.12 Fixed-Time Quanta Charts

Another way of visualizing application activity is a fixed-time quanta (FTQ) chart. While the call graph gives a very detailed profile of what code regions are most important for the application, they lack temporal information. The FTQ histogram gives a time-dependent profile of what the application is doing (Figure 3.19). This can be useful for observing the ratio of communication to computation. It can also give a sense of how “steady” the application is, i.e. if the application oscillates between heavy computation and heavy communication or if it keeps a constant ratio. In the simple example, Figure 3.19, we show the FTQ profile of a simple MPI test suite with random computation mixed in. In general, communication (MPI) dominates. However, there are a few compute-intensive and memory-intensive regions.

The FTQ visualization is activated by another input parameter

```
ftq = <fileroot>
```

where the `fileroot` parameter gives a unique prefix for the output files.

After running, two new files appear in the folder: `<fileroot>_app1.p` and `<fileroot>_app1.dat`. `plot_app1.p` is a Gnuplot script that generates the histogram as a postscript file.

```
your_project$ gnuplot plot_app1.p > output.ps
```

Gnuplot can be downloaded from <http://www.gnuplot.info> or installed via MacPorts. We recommend version 4.4, but at least 4.2 should be compatible.

The granularity of the chart is controlled by the `ftq_epoch` parameter in the input file. The above figure was collected with

```
ftq_epoch=5us
```

Events are accumulated into a single data point per “epoch.” If the timestamp is too small, too little data will be collected and the time interval won’t be large enough to give a meaningful picture. If the timestamp is too large, too many events will be grouped together into a single data point, losing temporal structure.

Using fully namespace parameters, this would be specified as:

```
node.os.ftq.fileroot=<fileroot>
node.os.ftq.epoch=5us
```

3.13 Network Statistics

Here we describe a few of the network statistics that can be collected and the basic mechanism for activating them. These statistics are usually collected on either the NIC, switch crossbar, or switch output buffers.

3.13.1 Message Size Histogram

To active a message size histogram on the NIC to determine the distribution of message sizes, the parameter file should include, for example:

```
node {
  nic {
    message_size_histogram {
      fileroot = histogram
      bin_size = 1
      logarithmic = true
    }
  }
}
```

The statistics are activated when the parameter reader sees the namespace `message_size_histogram`. In this case, we ask for a logarithmic distribution. The bin size here is in logarithmic units, i.e. group results in bins corresponding to an exponent range of size 1. An example generated for Nekbone with 1024 processors is in Figure 3.20.

3.13.2 Congestion Delay Histogram

A more involved example looks at congestion delays in the application. We want to generate a histogram showing the aggregate delay (relative to zero-congestion baseline) that a packet experiences going from source to destination. By default, packets do not carry fields for measuring congestion. Thus, the packet allocator must be changed. The NIC parameters now become:

```
node {
  nic {
    packet_allocator = delay_stats
    ejection {
      stats = delay_histogram
      delay_histogram {
        fileroot = delay
        bin_size = 0.5us
      }
    }
  }
}
```

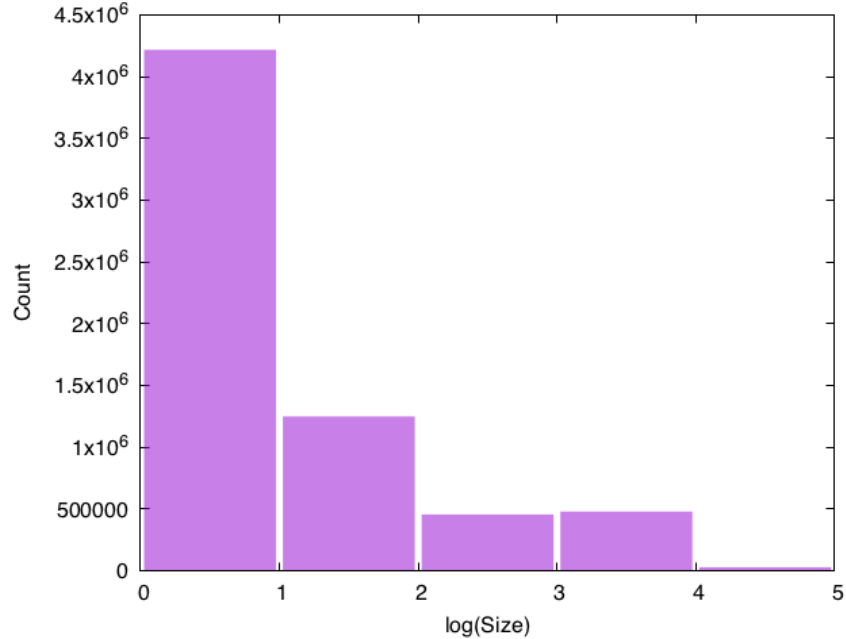


Figure 3.20: Logarithmic histogram of message sizes sent by Nekbone application

The delay histogram goes in the `ejection` namespace since we want to measure packet congestion after it leaves the network. The histogram is not logarithmic and we want to bin files in units of $0.5\mu s$. After running, two files will be generated: a gnuplot script `delay.p` and a corresponding data file `delay.dat`. After running

```
shell>gnuplot delay.p
```

a PNG file `delay.png` is generated. The generated histogram is shown in Figure 3.21. Delays, when occurring in this case, are usually on the order of a few μs .

A few more things are actually required to complete the input file. The input file above only tells the NIC to create a histogram. We still have to tell the crossbars and buffers on the network to accumulate the delays on the packet. Otherwise the NIC will just see a zero delay. The following extra parameters are required:

```
switch {
  xbar {
    stats = congestion_delay
  }
  output_buffer {
    stats = congestion_delay
  }
}
```

3.13.3 Congestion Spyplot and Multi-stats

Another way to look for congestion is to create a spyplot. Each row/column is a source/destination pair. In contrast to previous spyplots that show the amount of traffic, a congestion spyplot shows the amount of network congestion experienced




Figure 3.21: Histogram of message sizes sent by Nekbone application

sending between two points. Most of the same parameters from the delay histogram are required. However, we now want to collect both a spyplot and the histogram.

```
node {
  nic {
    packet_allocator = delay_stats
    ejection {
      stats = multi
      callbacks = congestion_spyplot delay_histogram
    }
    congestion_spyplot {
      fileroot = spyplot
      type = png
      normalization = 100000
    }
    delay_histogram {
      fileroot = delay
      bin_size = 0.5 us
    }
  }
}
```

We set the stats collector to `multi`. We then supply the list of desired stats to the `callbacks` parameter. Zooming into the Nekbone example, we can see congestion hotspots. For the most part, very little congestion appears in Figure 3.22. However (at least for the Torus topology used here), there are a few off-diagonal regions that show some congestion.

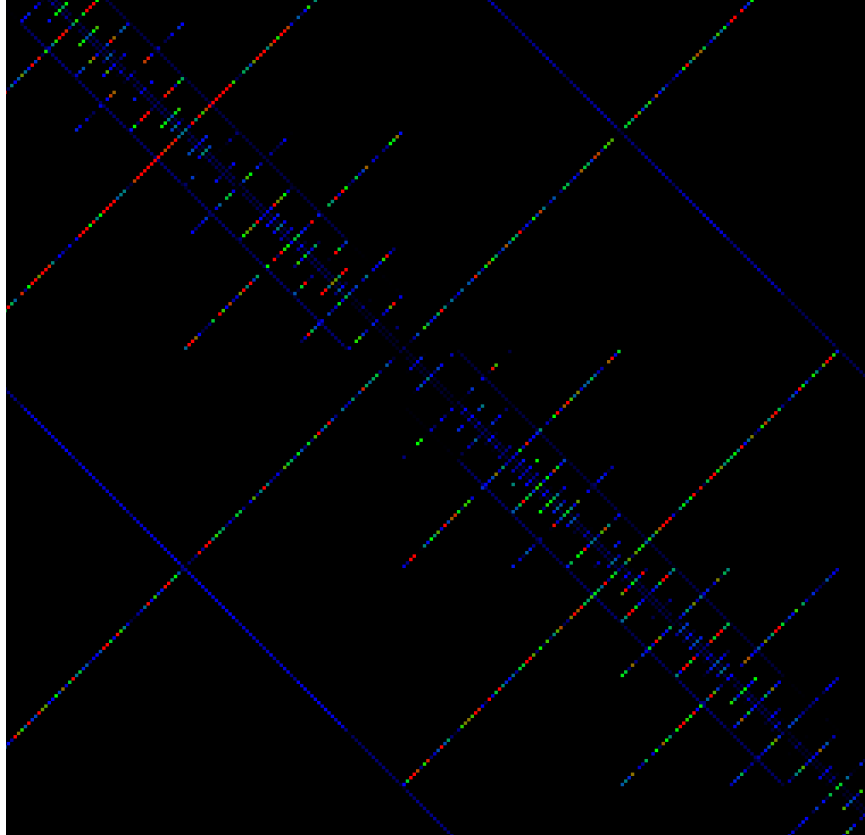


Figure 3.22: Spyplot showing congestion hotspots for certain source/destination pairs for Nekbone.

Chapter 4

Topologies

The torus topology is straightforward and easy to understand. Here we introduce the basics of other topologies within SST that are more complex and require extra documentation to configure properly. These are generally higher-radix or path-diverse topologies like fat tree, dragonfly, and flattened butterfly. As noted in 3.3, a more thorough and excellent discussions of these topologies is given in “High Performance Datacenter Networks” by Dennis Abts and John Kim.

4.1 Topology Query Utility

Understanding topology inputs and geometries can sometimes be challenging. SST/macro provides an executable for testing topology inputs and doing example coordinate computations. After making and installing, an executable `sstmac_top_info` will appear in the `bin` folder. The invocation of `sstmac_top_info` is exactly the same as the main `sstmac` executable. For the example parameter file named `machine.ini`:

```
topology.name = fattree
topology.geometry = 4 3
```

we run

```
bin> sstmac_top_info -f machine.ini
```

which produces the output

```
Number of nodes:      81
Number of leaf switches: 27
Number of switches:   94
```

detailing the produced geometry. Here the fat tree has a total of 94 switches, 27 of which are “leaf” switches directly connected to compute nodes. The output is followed by the prompt

NextInput:

One can either enter a single number (switch ID) or set of coordinates. If given a switch ID, the coordinates are computed. If coordinates are given, the switch ID is computed.

```
NextInput: 32
Switch ID maps to coordinates [ 2 0 1 2 ]
NextInput: 2 0 1 2
Coordinates map to switch ID 32
```

The program is just exited with Ctrl-C. The meaning of the above coordinates is detail below for fat tree (Section 4.4).

4.2 Torus

The torus is the simplest topology and fairly easy to understand. We have already discussed basic indexing and allocation as well as routing. More complicated allocation schemes with greater fine-grained control can be used such as the coordinate allocation scheme (see hypercube below for examples) and the node ID allocation scheme (see fat tree below for examples). More complicated Valiant and UGAL routing schemes are shown below for hypercube and dragonfly, but apply equally well to torus.

For torus we illustrate here the Cartesian allocation for generating regular Cartesian subsets. For this, the input file would look like

```
topology {
  name = torus
  geometry = 4 4 4
}
node {
  appl {
    launch_cmd = aprun -n 8
    indexing = block
    allocation = cartesian
    cart_sizes = 2 2 2
    cart_offsets = 0 0 0
  }
}
```

This allocates a 3D torus of size 4x4x4. Suppose we want to allocate all 8 MPI ranks in a single octant? We can place them all in a 2x2x2 3D sub-torus by specifying the size of the subblock (**cart_sizes**) and which octant (**cart_offsets**). This applies equally well to higher dimensional analogs. This is particularly useful for allocation on Blue Gene machines which always maintain contiguous allocations on a subset of nodes.

This allocation is slightly more complicated if we have multiple nodes per switch. Even though we have a 3D torus, we treat the geometry as a 4D coordinate space with the 4th dimension referring to nodes connected to the same switch, i.e. if two nodes have the 4D coordinates [1 2 3 0] and [1 2 3 1] they are both connected to the same switch. Consider the example below:

```
topology {
  name = torus
  geometry = 4 4 4
  concentration = 2
}
appl {
  launch_cmd = aprun -n 8
  indexing = block
  allocation = cartesian
  cart_sizes = 2 2 1 2
  cart_offsets = 0 0 0 0
}
```

We allocate a set of switches across an XY plane (2 in X, 2 in Y, 1 in Z for a single plane). The last entry in **cart_sizes** indicates that both nodes on each switch should be used.

4.3 Hypercube

Although never used at scale in a production system, the generalized hypercube is an important topology to understand, particularly for flattened butterfly and dragonfly. The (k,n) generalized hypercube is geometrically an N -dimensional torus with each dimension having size k (although dimension sizes need not be equal). Here we show a $(4,2)$ generalized hypercube (Figure 4.1). This would be specified in SST as:

```
topology.name = hypercube
topology.geometry = 4 4
```

indicating size 4 in two dimensions.

While a torus only has nearest-neighbor connections, a hypercube has full connectivity within a row and column (Figure 4.1). Any switches in the same row or same column can send packets with only a single hop.

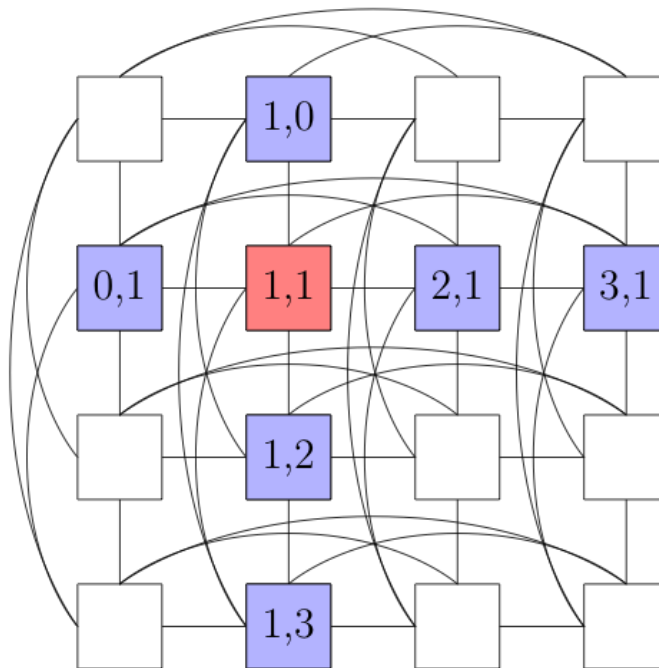


Figure 4.1: Hypercube with links and connections within a row/column

This extra connectivity leads to greater path diversity and higher radix switches. The cost tradeoff is that each link has lower bandwidth than a torus. Whereas a torus has a few fat links connecting switches, a hypercube has many thin links. A hypercube can have more dimensions and be asymmetric, e.g.

```
topology.name = hypercube
topology.geometry = 4 5 6
```

where now we have full connections within horizontal rows, horizontal columns, and vertical columns. Here each switch has radix 12 (3 connections in X, 4 connections in Y, 5 connections in Z).

4.3.1 Allocation and indexing

A hypercube has the same coordinate system as a torus. For example, to create a very specific, irregular allocation on a hypercube:

```
node {  
  appl {  
    launch_cmd = aprun -n 5  
    indexing = coordinate  
    allocation = coordinate  
    coordinate_file = coords.txt  
  }  
}
```

and then a coordinate file named `coords.txt`

```
5 2  
0 0  
0 1  
1 1  
2 0  
3 3
```

The first line indicates 5 entries each with 2 coordinates. Each line then defines where MPI ranks 0-4 will be placed

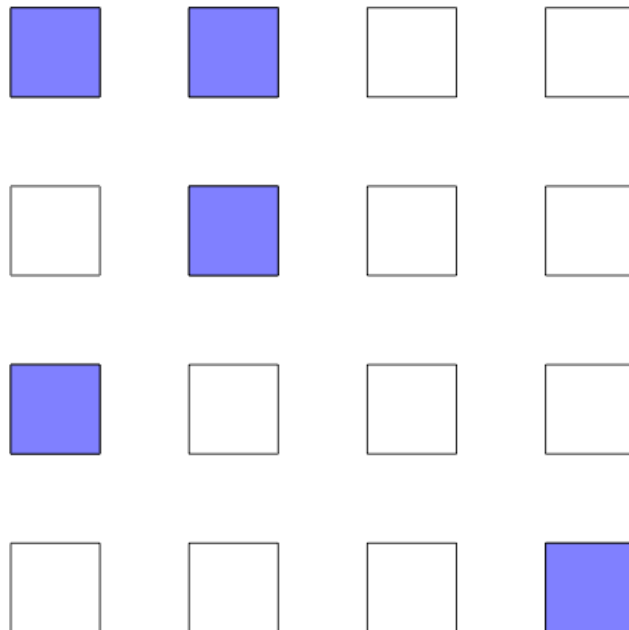


Figure 4.2: Hypercube allocation for given set of coordinates

4.3.2 Routing

Hypercubes allow very path-diverse routing because of its extra connections. In the case of minimal routing (Figure 4.3), two different minimal paths from blue to red are shown. While dimension order routing would rigorously go X then Y, you can still route minimally over two paths either randomly selecting to balance load or routing based on congestion.




Figure 4.3: Minimal routing within a hypercube showing path diversity. Packet travels from blue to red, passing through green intermediate switches.

To fully maximize path diversity on adversarial traffic patterns, though, path-diverse topologies can benefit from Valiant routing. Here, rather than directly routing to the final destination, packets first route to random intermediate switches on a minimal path. Then they route again from the intermediate switch to the final destination also on a minimal path (Figure 4.4). Although it increases the hop count and therefore the point-to-point latency, it utilizes more paths and therefore increases the effective point-to-point bandwidth.

4.4 Fat Tree

Within SST, a fat tree is defined by the following parameters:

```
topology.name = fattree
topology.geometry = 4 2
```

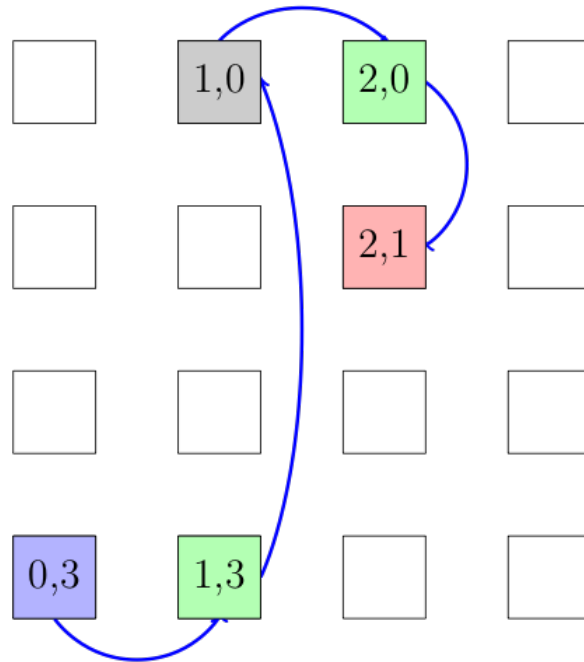


Figure 4.4: Valiant routing within a hypercube. Packet travels from blue to red via a random intermediate destination shown in gray. Additional intermediate switches are shown in green.

The first number, 4, indicates the number of levels in the fat tree. The second number, 2, indicates the radix or branching factor of the tree. The number of compute nodes in this topology is $2^4 = 16$. This is illustrated conceptually in Figure 4.5. The color coding will become clear later. We note this is somewhat confusing since the fat tree appears to have 5 levels. Here the topology is defined by the number of levels containing switches or the number of branches. This is done for a very specific reason. At the final level, you may wish to have a different branching fraction for the compute nodes, e.g.

```
topology.concentration = 1
```

This loads the injection bandwidth for the compute node dedicating its own injection switch. If the parameter `network_nodes_per_switch` is omitted, it defaults to the fat tree radix. This case is shown in Figure 4.5 where there are two nodes injecting to the same switch. Higher radix fat trees can be specified, e.g.

```
topology.name = fattree
topology.geometry = 3 4
```

which would have $4^3 = 64$ compute nodes.

In reality, it is not practical to implement a fat tree exactly as shown in Figure 4.5. One would need to buy many non-standard, high capacity switches for the higher levels in the fat-tree. The simple model is available by specifying `simple_fattree` as the topology, and SST will construct special large switches at higher levels. The best practical implementation employs all uniform, commodity switches (Figure 4.6). The fat tree is “virtual” with several commodity switches grouped together




Figure 4.5: Abstract, conceptual picture of Fat Tree topology

to simulate a heavy-weight, high capacity switch at higher levels of the fat tree. The connection between the physical implementation and the conceptual fat tree can easily be seen by the color coding. For example, the second row contains eight switches, but only two virtual switches. Each virtual switch is composed of four commodity switches.




Figure 4.6: Physical implementation of Fat Tree with commodity switches showing ID numbering

Within SST, each switch is assigned a unique ID, starting from zero in the bottom row and proceeding through the top level. In addition, each compute node is also assigned a unique ID from 0 to 15. The switches can also be defined by a set of coordinates. While the choice of coordinate system for a 3D torus is obvious, the coordinate system for the fat tree is less clear. In SST, we define a 2D mesh coordinate system for the row (level) and column of the switch.

4.4.1 Allocation and indexing

The numbering of compute nodes is shown in Figure 4.6. Consider the case

```
node.app1.launch_cmd = aprun -n 4 -N 1
```

which launches four processes all on distinct nodes. In the simplest allocation and indexing scheme (first available), processes would be placed in order on 0,1,2,3. An alternative allocation/indexing scheme uses the Node ID allocator.

```
node {
  app1 {
    allocation = node_id
    indexing = node_id
    node_id_file = nodes.txt
  }
}
```

Here `nodes.txt` would contain the number of nodes on the first line, followed by the list of Node IDs, in order, of where to place MPI ranks. For the file

```
4
0
4
8
12
```

Four MPI ranks would be placed in spatially distant parts of the machine.

If indexing differs from allocation (usually because there are multiple MPI ranks per node), both an allocation and an indexing file are needed. Suppose we have:

```
node.app1.launch_cmd = aprun -n 4 -N 2
```

We then need:

```
node {
  app1 {
    allocation = node_id
    indexing = node_id
    node_id_allocation_file = alloc.txt
    node_id_indexing_file = index.txt
  }
}
```

where the contents of `alloc.txt` are, e.g.

```
2
0
1
```

choosing nodes 0 and 1 in the allocation and then `index.txt` would be, e.g.

```
4
0
1
0
1
```

which round-robin assigns rank 0 to node, rank 1 to node 1, rank 2 to node 0, and so on.

4.4.2 Routing

Fat tree routing is actually straightforward, but can employ path diversity. Suppose you are routing from Node 0 to Node 2 (Figure 4.6). At the first stage, you have no choice. You must route to Switch 1. At the second stage, you can either route to Switch 8 or Switch 9. Suppose you branch to Switch 9. At this point, you are done moving up. The packet now proceeds down the fat-tree. On the downward routing, there is no path diversity. Only a single, minimal route exists to the destination node. In the simplest case, Switch 1 alternates between selecting Switch 8 and Switch 9 to distribute load. In a more complicated scheme, Switch 1 could adaptively route selecting either Switch 8 or Switch 9 based on congestion information.

4.5 Dragonfly

As bandwidth per pin increases, arguments can be made that optimal topologies should be higher radix. A 3D torus is on the low-radix extreme while a hypercube is a high-radix extreme. Unfortunately a hypercube topology is not scalable and the radix quickly becomes too high to efficiently implement. A dragonfly is sometimes viewed as a generalization of flattened butterfly and hypercube topologies with “virtual” switches of very high radix, not dissimilar from the fat-tree implementation with many physical commodity switches composing a single virtual switch. The dragonfly topology (Figure 4.7) is actually quite simple. Small groups are connected as a generalized hypercube with full connectivity within a row or column. Intergroup connections (global links) provide pathways for hopping between groups. A dragonfly is usually understood through three parameters:

- p : number of nodes connected to each router
- a : number of routers in a group
- h : number of global links that each switch has

For simplicity, only three example global links are show for clarity in the picture. For the Cray X630, $a = 96$, $h = 10$, and $p = 4$ so that each router is connected to many other ($h = 10$) groups. The caveat is that in many implementations global links are grouped together for $h = 2$ or 3 fat global links. These demonstrate well-balanced ratios. In general, scaling out a dragonfly should not increase the size of a group, only the number of groups.

4.5.1 Allocation and indexing

The dragonfly coordinate system is essentially the same as a 3D torus. The group 2D hypercube layout defines X and Y coordinates. The group number defines a Z or G coordinate. Thus the topology in Figure 4.7 would be specified as

```
topology.name = dragonfly
topology.geometry = 3 3 3
```

for groups of size 3×3 with a total of 3 groups. To complete the specification, the number of global links (h) for each router must be given

```
topology.group_connections = 10
```




Figure 4.7: Schematic of dragonfly with three groups showing hypercube intragroup links and high bandwidth intergroup global links

4.5.2 Routing

It is important to understand the distinction between link bandwidth, channel bandwidth, and pin bandwidth. All topologies have the same pin bandwidth and channel bandwidth (assuming they use the same technology). Each router in a topology is constrained to have the same number of channels (called radix, usually about $k = 64$). The number of channels per link changes dramatically from topology to topology. Low radix topologies like 3D torus can allocate more channels per link, giving higher bandwidth between adjacent routers. Dragonfly is higher radix, having many more connections but having lower bandwidth between adjacent routers. While minimal routing is often sufficient on torus topologies because of the high link bandwidth, dragonfly will exhibit very poor performance with minimal routing. To effectively utilize all the available bandwidth, packets should have a high amount of path diversity. Packets sent between two routers should take as many different paths as possible to maximize the effective bandwidth point-to-point.

Minimal routing itself has a few complications (Figure 4.8). Each router only has a few global links. Thus, traveling from

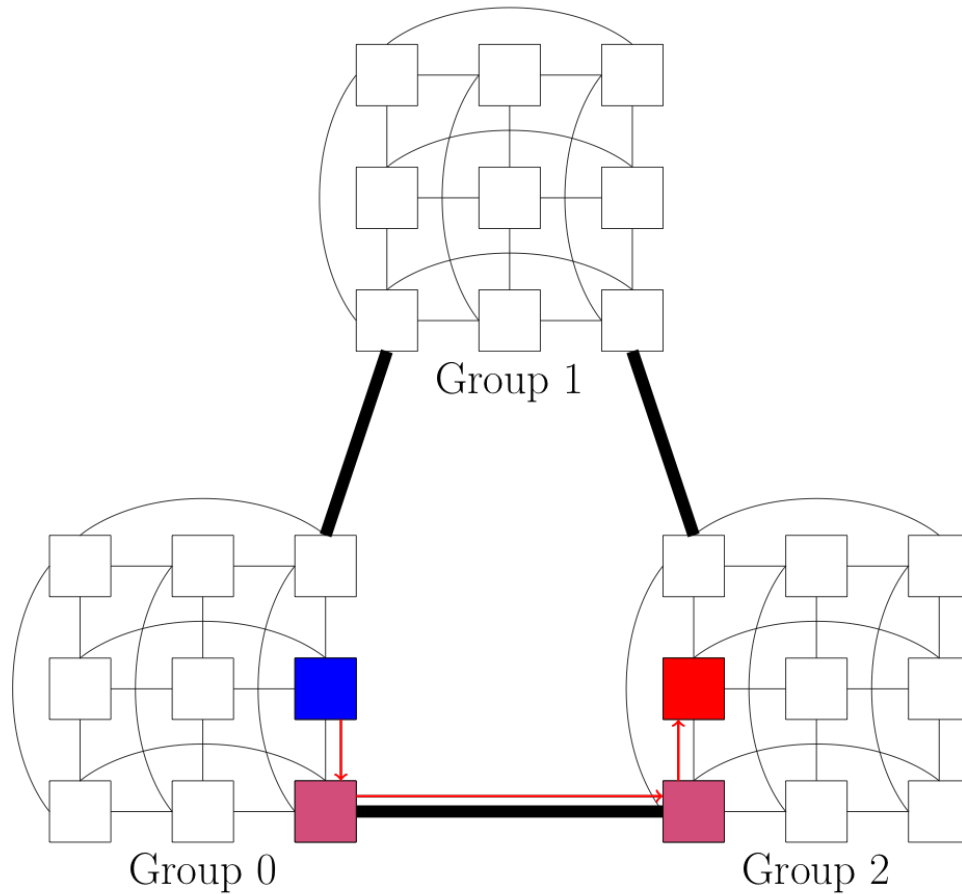


Figure 4.8: Schematic of dragonfly showing minimal route. Traveling between groups requires routing to the correct global link, hopping the global link, then routing within a group to the correct final node.

e.g. the blue router at $X=3, Y=2, G=0$ to the red router at $X=1, Y=2, G=2$, there is no direct link between the routers. Furthermore, there is no direct link between Groups 0 and 2. Thus packets must route through the purple intermediate nodes. First, the packet hops to $X=3, Y=3, G=0$. This router has a global link to Group 2, allowing the packet to hop to the next intermediate router at $X=1, Y=3, G=2$. Finally, the minimal route completes by hopping within Group 2 to the final destination.

To improve on minimal routing, global routing strategies are required (global routing is distinguished here from adaptive routing). Global essentially means “not minimal” and spreads packets along many different paths. The simplest global routing strategy is Valiant routing, which falls in the global, oblivious category (Figure 4.9). Oblivious simply means packets are scattered randomly without measuring congestion. In Valiant routing, each packet does the following:

- Pick a random intermediate node
- Route minimally to random node




Figure 4.9: Schematic of dragonfly showing Valiant route. Traveling between groups requires routing to a random intermediate node, then routing minimally to the final destination.

- Route minimally from random node to destination node

This is somewhat counterintuitive at first. Rather than go directly to the destination node, packets go out of their way to a random node, shown in Figure 4.9 as the yellow router. Thus, routing from the blue router in Group 0 to the red router in Group 2 first follows the minimal path (green routers) to the randomly selected yellow router in Group 1. From there, a second minimal path is taken through the orange routers to the final destination. In cases with high congestion or even for large messages on a quiet network, this actually improves performance. If a point-to-point message is composed of ten packets, all ten packets will follow different paths to the final destination. This essentially multiplies the maximum bandwidth by a factor of ten. Valiant routing can be specified as

```
router = valiant
```

In contrast, UGAL routing is a global, adaptive strategy, making decisions based on congestion. Because Valiant is oblivious,

it often sends too many packets to far away random nodes. Following a Valiant path is only relevant when enough packets fill up router queues, creating congestion. UGAL does the following steps:

- Start routing minimally
- On each step, check congestion (buffer queue depth)
- If congestion is too heavy, switch to Valiant and re-route to random intermediate node. Otherwise stay on minimal path.

UGAL packets stay on a minimal path until congestion forces them to use a Valiant strategy. This routing can be specified as:

```
router = ugal
```

Chapter 5

External Applications and Skeletonization

5.1 Basic Application porting

There are three parts to successfully taking a C++ code and turning it into a running application.

- **Redirected linkage:** Rather than linking to MPI, pThreads, or other parallel libraries (or even calling `hostname`), these functions must be redirected to SST/macro rather than calling the native libraries on the machine running the simulator. You get all redirected linkage for free by using the SST compiler wrappers `sst++` and `sstcc` installed in the `bin` folder.
- **Skeletonization:** While SST/macro can run in emulation mode, executing your entire application exactly, this is not scalable. To simulate at scale (i.e. 1K or more MPI ranks) you must strip down or “skeletonize” the application to the minimal amount of computation. The energy and time cost of expensive compute kernels are then simulated via models rather than explicitly executed.
- **Process encapsulation:** Each virtual process being simulated is not an actual physical process. It is instead modeled as a lightweight user-space thread. This means each virtual process has its own stack and register variables, but not its own data segment (global variables). Virtual processes share the same address space and the same global variables. A Beta version of the auto-skeletonizing clang-based SST compiler is available with the 7.X releases. If the Beta is not stable with your application, manual refactoring may be necessary if you have global variables.

5.1.1 Loading external skeletons with the standalone core

You should always write your own skeleton applications in an external folder rather than integrating directly into the `sstmac` executable. You have two options for running your skeleton app within the simulation:

- Create a new executable with your skeleton app built-in
- Create a library `libX.so` and have `sstmac` load it at runtime

New Executable

If you follow the example in the `skeletons/sendrecv` folder, the Makefile shows how to generate a new executable. If you are using `sst++`, it will automatically link your executable with the right SST libraries. When you run your created executable, it will spawn an SST/macro simulation of your app. If your application is named `runapp`, you would run it exactly like the `sstmac`:

```
./runapp -f parameters.ini
```

where your new executable replaces `sstmac`.

Dynamically Loaded Library

If you follow the example in the `skeletons/sendrecv_sst_so` folder, the Makefile shows how to generate a skeleton application as an external library. Instead of replacing the `sstmac` executable, you tell the `sstmac` to load the library in the parameter file.

```
external_libs = libsendrecv.so
```

The library `libsendrecv.so` can be generated with standard compiler flags for shared libraries by using the `sst++` compiler. Once the `external_libs` parameter is set in the file, you can run:

```
sstmac -f parameters.ini
```

with the standard executable `sstmac`.

5.1.2 Loading external skeletons with the integrated core

While the main `libmacro.so` provides the bulk of SST/macro functionality, users may wish to compile and run external skeletons. This is basically the same workflow as shown above for dynamically loaded libraries. When the executable is generated, SST/macro will also generate a `libX.so` containing all the element info. This can then be used for SST simulations. The `default.py` script used by `pysstmac` must also be edited. The top lines was previously

```
import sst.macro
```

This only loads the main components, not the external skeleton. You must add

```
import sst.X
```

where X is the name of your skeleton. This causes the core to also load the shared library corresponding to your external skeleton app. If using the SST compiler wrappers, the ELI block and `.so` file will actually be generated automatically. As shown in Section 2.5, the generate shared library files can be added as the first parameters to the `pysstmac` script.

```
shell>sst-macro/skeletons/sendrecv> pysstmac librunsstmac.so -f parameters.ini
```

5.2 Auto-skeletonization with Clang (Beta)

The build of the Clang toolchain is described in Section 2.4. This enables a source-to-source translation capability in the `sst++` compiler that can auto-skeletonize computation and fix global variable references. Some of this can be accomplished automatically (global variables), but most of it (removing computation and memory allocations) must occur through pragmas. A good example of skeletonization can be found in the `lulesh2.0.3` example in the `skeletons` folder. Most of the available SST pragmas are used there. Pragmas are preferred since they allow switching easily back and forth between skeleton and full applications. This allows much easier validation of the simulation. The section here briefly introduces the SST pragma language. A complete tutorial on all available pragmas is given in Chapter 6.

5.2.1 Redirecting Main

Your application's `main` has to have its symbols changed. The simulator itself takes over `main`. SST/macro therefore has to capture the function pointer in your code and associate it with a string name for the input file. This is automatically accomplished by defining the macro `sstmac_app_name` either in your code or through a `-D=` build flag to the name of your application (unquoted!). The value of the macro will become the string name used for launching the application via `node.app1.name=X`. Even without Clang, this works for C++. For C, Clang source-to-source is required.

5.2.2 Memory Allocations

To deactivate memory allocations in C code that uses `malloc`, use:

```
1 #pragma sst malloc
2 void* ptr = malloc(...)
```

prior to any memory allocations that should be deactivated during skeleton runs, but active during real runs.

Similarly, for C++ we have

```
1 #pragma sst new
2 int* ptr = new int[...]
```

5.2.3 Computation

In general, the SST compiler captures all `#pragma omp parallel` statements. It then analyzes the for-loop or code block and attempts to derive a computational model for it. The computational models are quite simple (skeleton apps!), based simply on the number of flops executed and the number of bytes read (written) from memory. Consider the example:

```
1 double A[N], B[N];
2 #pragma omp parallel for
3 for (int i=0; i < N; ++i){
4     A[i] = alpha*A[i] + B[i];
5 }
```

The SST compiler deduces $16N$ bytes read, $8N$ bytes written, and $16N$ flops (or $8N$ if fused-multiplies are enabled). Based on processor speed and memory speed, it then estimates how long the kernel will take without actually executing the loop. If not wanting to use OpenMP in the code, `#pragma sst compute` can be used instead of `#pragma omp parallel`.

5.2.4 Special Pragmas

Many special cases can arise that break skeletonization. This is often not a limit of the SST compiler, but rather a fundamental limitation in the static analysis of the code. This most often arises due to nested loops. Consider the example:

```
1 #pragma omp parallel for
2 for (int i=0; i < N; ++i){
3     int nElems = nElemLookup[i];
4     for (int e=0; e < nElems; ++e){
5     }
6 }
```

Auto-skeletonization will fail. The skeletonization converts the outer loop into a single call to an SST compute model. However, the inner loop can vary depending on the index. This data-dependency breaks the static analysis. To fix this, a hint must be given to SST as to what the “average” inner loop size is. For example, it may loop nodes in a mesh. In this case, it may almost always be 8.

```
1 #pragma omp parallel for
2 for (int i=0; i < N; ++i){
3     int nElems = nElemLookup[i];
4     #sst replace nElems 8
5     for (int e=0; e < nElems; ++e){
6     }
7 }
```

This hint allows SST to skeletonize the inner loop and “guess” at the data dependency.

5.2.5 Skeletonization Issues

Skeletonization challenges fall into three main categories:

- *Data structures* - Memory is a precious commodity when running large simulations, so get rid of every memory allocation you can.
- *Loops* - Usually the main brunt of CPU time, so get rid of any loops that don’t contain MPI calls or calculate variables needed in MPI calls.
- *Communication buffers* - While you can pass in real buffers with data to SST/macro MPI calls and they will work like normal, it is relatively expensive. If they’re not needed, get rid of them.

The main issue that arises during skeletonization is data-dependent communication. In many cases, it will seem like you can’t remove computation or memory allocation because MPI calls depend somehow on that data. The following are some examples of how we deal with those:

- *Loop convergence* - In some algorithms, the number of times you iterate through the main loop depends on an error converging to near zero, or some other converging mechanism. This basically means you can’t take out anything at all, because the final result of the computation dictates the number of loops. In this case, we usually set the number of main loop iterations to a fixed number.

- *Particle migration* - Some codes have a particle-in-cell structure, where the spatial domain is decomposed among processes, and particles or elements are distributed among them, and forces between particles are calculated. When a particle moves to another domain/process, how many particles migrate and how far depends on the actual computed forces. However, in the skeleton, we are not actually computing the forces - only estimated how long the force computation took. If all we need to know is that this migration/communication happens sometimes, then we can just make it happen every so many iterations, or even sample from a probability distribution.
- *AMR* - Some applications, like adaptive mesh refinement (AMR), exhibit communication that is entirely dependent on the computation. In this case, skeletonization again depends on making approximations or probability estimates of where and when box refinement occurs without actually computing everything.

For applications with heavy dynamic data dependence, we have the following strategies:

- *Traces* - revert to DUMPI traces, where you will be limited by existing machine size. Trace extrapolation is also an option here.
- *Synthetic* - It may be possible to replace communication with randomly-generated data and decisions, which emulate how the original application worked. This occurs in the CoMD skeleton.
- *Hybrid* - It is possible to construct meta-traces that describe the problem from a real run, and read them into SST/macro to reconstruct the communication that happens. This occurs in the `boxml` applications.

5.3 Process Encapsulation

As mentioned above, virtual processes are not real, physical processes inside the OS. They are explicitly managed user-space threads with a private stack, but without a private set of global variables. When porting an application to SST/macro, global variables used in C programs will not be mapped to separate memory addresses causing incorrect execution or even segmentation faults. If you have avoided global variables, there is no major issue. If you have read-only global variables with the same value on each machine, there is still no issue. If you have mutable global variables, you should use the `sst++` clang-based compiler wrappers to auto-refactor your code (Section 5.2). This feature is current labeled Beta, but is stable for numerous tests and will be fully supported for release 7.1.

Chapter 6

Clang Source-to-Source Auto-Skeletonization via Pragmas

There are three main examples of auto-skeletonization with pragmas in the SST/macro source code in the `skeletons` directory. These applications are Lulesh, HPCG, and CoMD. The auto-skeletonizing compiler is designed to do three main things:

- Redirect global variable accesses to thread-specific values
- Turn off large memory allocations that would prevent scalable simulation
- Estimate time of compute-intensive kernels instead of executing them

6.1 Pragma Overview

6.1.1 Compiler workflow

The source-to-source compiler operates on a pre-processed source file. The source code transformation generates a temporary source file. This temporary source file is then compiled into the target object file. Global variables require static registration of C++ variables. Here another temporary C++ source file (even if the original file is C) is generated that has all static global variable registrations. The corresponding object file is merged with the original object file, creating a complete SST/macro object file with the transformed code and C++ static registrations. This workflow is shown in Figure 6.1.

6.1.2 Compiler Environment Variables

SSTMAC_SRC2SRC: Default 1

If set to zero, deactivates the source-to-source transformation. The compiler wrapper will then compile the code into the simulator, but will not redirect any global variable accesses or perform any skeletonization.




Figure 6.1: Source-to-source transformation workflow for SST compiler. For C source files, g++ can be swapped with gcc. The choice of underlying compiler is actually arbitrary and can be clang, gcc, icc, etc.

SSTMAC_SKELETONIZE: Default 0

If set to zero, deactivates skeletonization. This does not deactivate global variable redirection. Thus, with SSTMAC_SRC2SRC=1 and SSTMAC_SKELETONIZE=0, SST/macro will act as an MPI emulator executing a full code but with global variables refactored

to maintain correctness.

SSTMAC_HEADERS: No default

The compiler wrapper will only redirect global variables that it knows should definitely be modified. All global variables found in source files will be redirected. `extern` global variables found in header files are more difficult. Certain system global variables like `stderr` should not be modified and so are left as global variable constants. By default, global variables in a header file area NOT redirected unless explicitly specified in a header configuration file. The variable `SSTMAC_HEADERS` should give the full path of a file containing the list of header files. Header file paths in the file should be one per line and should be the full path, not a relative path.

SSTMAC_DELETE_TEMPS: Default 1

If non-zero, the compiler cleans up all temporary files. If you wish to keep temporary files to view them for debugging, set to zero. All temporary, intermediate source files will not be deleted at the end of compilation.

6.2 Basic Replacement Pragmas

When skeletonization is active (see `SSTMAC_SKELETONIZE`, these pragmas will cause replacements in the original source code. Pragmas apply to the next statement in the source code. For compound statements such as a for-loop with a multi-statement body, the pragma applies to the entire for-loop.

6.2.1 `pragma sst delete: no arguments`

This deletes the next statement from the source code. If the statement declares a variable that is used later in the code, this will cause a compile error. Consider an example from the Lulesh source code.

```
1 #pragma sst delete
2 testnorms_data.values[i] = normr/normr0;
```

In the skeleton, the residual is not actually computed and the `testnorms_data` array is not actually allocated. Thus this statement should be deleted and not actually executed in the skeleton.

6.2.2 `pragma sst replace [to_replace:string] [new_text:C++ expression]`

This applies a string replace to a variable or function call in the next statement. Consider an example from Lulesh.

```
1 #pragma sst replace volo 1
2 deltatime() = (Real_t(.5)*cbrt(volo(0)))/sqrt(Real_t(2.0)*einit);
```

The function call `volo(0)` is not valid in the skeleton since volumes are not actually computed. Here we simply estimate that all cells have unit volume replacing `volo(0)` with 1.

6.2.3 `pragma sst init [new_value:string]`

Can only apply to a binary equal operator assigning a value. The pragma changes the right-hand side to use the given new value. For example, in Lulesh:

```
1 #pragma sst init nullptr
2   destAddr = &domain.commDataSend[pmsg * maxPlaneComm] ;
```

The send buffer `domain.commDataSend` is not allocated in the skeleton and thus is not valid to use. The pragma causes the skeleton to simply assign `destAddr` to `nullptr`.

6.2.4 `pragma sst return [new_value:C++ expression]`

Equivalent to `pragma sst init`. Replaces the target of a return statement with the given expression. This produces a compiler error if applied to anything but a return statement.

6.2.5 `pragma sst keep`

During the skeletonization process, some transformations occur automatically even without pragmas. For example, all MPI calls have their input buffers converted to null pointers to indicate a simulated MPI call. If the MPI call should be emulated and actually deliver a real payload, the MPI call must be explicitly marked with `pragma keep`. An example can be found in the HPCG skeleton:

```
1 #pragma sst keep
2   MPI_Allreduce(&localNumberOfNonzeros, &totalNumberOfNonzeros, ...)
```

The actual allreduce operation is carried out, summing the the local number into the total number of non zeroes.

6.2.6 `pragma sst keep_if [condition:C++ bool expression]`

More control over whether a transformation is skipped can be provided by the `keep_if` program. An example can be found in CoMD.

```
1 #pragma sst keep_if count < 16
2   MPI_Allreduce(sendBuf, recvBuf, count, MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
```

Any small allreduce operations are kept. Any allreduce operations larger than a given cutoff are simulated without emulating the actual buffer operations.

6.2.7 `pragma sst empty`

Applied to functions. The function prototype is kept, but an empty body is inserted instead of the actual code. This can be useful for deleting large blocks of computation inside a function.

6.2.8 `pragma sst branch_predict [condition:C++ expression]`

The branch prediction pragma can be applied in two different contexts. We will revisit the pragma in the context of compute skeletonization below. The branch prediction pragma should only be applied to an if-statement. Much like the replace pragmas, it swaps out the given if condition with a new expression.

The branch prediction pragmas become necessary when skeletonizing. Certain values may not be computed or certain variables marked null. If these values are then used in an if predicate, the skeletonizing compiler cannot deduce the correct behavior for the application. When an ambiguous predicate is found, the compiler will usually print a warning and just assume the predicate is false. Predicates are often almost always true or almost always false. Thus most uses of this pragma will simply supply `true` or `false` as the replacement. However, any arbitrary C++ boolean expression can be given as the replacement. The new predicate expression (like the original), must not involve any null variables.

6.3 Memory Allocation Pragmas

6.3.1 `pragma sst malloc`

Applied to any statement in which the right-hand side is a malloc. This sets the left-hand side to `NULL`. This is critical for turning off large memory allocations on data structures not required for control-flow.

6.3.2 `pragma sst new`

Applied to any statement in which the right-hand side a C++ operator `new`. This sets the left-hand side to `nullptr`. This is critical for turning off large memory allocations on data structures not required for control-flow.

6.4 Data-Driven Type Pragmas

6.4.1 `pragma sst null_variable`

This applies to variable declarations. If pragma is not applied to a declaration, a compiler error is given. A null variable is one in which all operations involving the variable should be deleted. This usually applies to large data arrays that should never be allocated and therefore never dereferenced. An example can be seen in CoMD:

```
1 #pragma sst null_variable
2 int* nAtoms;           //!< total number of atoms in each box
```

The array is not allocated and all statements operating on the array are deleted.

This pragma is much more powerful than simply using `pragma sst new`. `pragma sst new` simply turns off a given memory allocation setting it a null value. If the array is dereferenced later in code, this causes a segmentation fault. By marking a declaration as null, the compiler can flag where segmentation faults would occur when running the skeleton.

In most cases, all operations involving the null variable are deleted. However, there may be cases where the compiler may decide deleting an operation cannot be done automatically since it may affect control flow, e.g., if the variable is used inside an if-statement. When this occurs, a compiler error is throw flagging where the ambiguity occurs. Another pragma must then be applied to that statement to tell the compiler how to proceed.

6.4.2 `pragma sst null_type` [type alias] [list allowed functions]

This applies to C++ class variable declarations. If `pragma` is not applied to a declaration, a compiler error is given. This essentially works the same way as `null_variable`, but allows certain member functions to be kept instead of deleted. Consider an example from Lulesh:

```
1 #pragma sst null_type sstmac::vector size resize empty
2   std::vector<Real_t> m_x ; /* coordinates */
```

Here we wish to indicate the vector is “null” and should not actually allocate memory or allow array accesses. However, we still wish to track the vector size and whether it is empty. The first argument to the `pragma` is a new type name that implements the “alias” functionality. For `std::vector`, SST/macro automatically provides the alias. For illustration, that code is reproduced here:

```
1 namespace sstmac {
2   class vector {
3   public:
4     void resize(unsigned long sz){
5       size_ = sz;
6     }
7
8     unsigned long size() const {
9       return size_;
10    }
11
12    template <class... Args>
13    void push_back(Args... args){
14      ++size_;
15    }
16
17    template <class... Args>
18    void emplace_back(Args... args){
19      ++size_;
20    }
21
22    bool empty() const {
23      return size_ == 0;
24    }
25
26    private:
27      unsigned long size_;
28  };
29 }
```

This empty vector class allows the type to track its size, but not actually hold any data. All places in the Lulesh code that an `std::vector` is used are substituted with the new type.

The remaining arguments to the `pragma` are the list of functions we wish to mark as valid. In this case, even though the alias vector class provides more functions, we only allow `size`, `resize`, and `empty` to be called.

6.5 Compute Pragas

6.5.1 `pragma sst compute` and `pragma omp parallel`

Compute-intensive should not be executed natively. Instead, a compute model should be used to estimate the elapsed time. Compute model replacements are automatically triggered by any OpenMP parallel pragmas. The corresponding block or

for-loop is not executed, instead calling out to a compute model to estimate time. Currently, compute modeling is done via a very basic static analysis. The static analysis attempts to count the number of integer and floating point operations. It also estimates the number of memory reads and writes. These four counters are passed to a coarse-grained processor model for time estimates. For more details, see `sstmac_compute_detailed` in the source code. Numerous examples can be found in the Lulesh, HPCG, and CoMD skeleton applications.

6.5.2 `pragma sst loop_count` [integer: C++ expression]

If the `sst compute` or `omp parallel` pragma are applied to an outer loop with one or more inner loops, the compute model static analysis might fail. This occurs when the inner loop control flow depends on the actual execution. Any variables declared *inside* the compute block are not valid to use in the compute estimate since they will be skeletonized and deleted. Only variables in scope at the beginning of the outer loop are safe to use in compute modeling.

When the static analysis fails, a corresponding compiler error is thrown. This usually requires giving a loop count hint. Consider the example from HPCG:

```
1  #pragma omp parallel for
2  for (local_int_t i=0; i< localNumberOfRows; i++) {
3      int cur_nnz = nonzerosInRow[i];
4      #pragma sst loop_count 27
5      for (int j=0; j<cur_nnz; j++) mtxIndL[i][j] = mtxIndG[i][j];
6  }
```

The static analysis fails on `cur_nnz`. However, that value is almost always 27. Thus we can safely tell the compiler to just assume the loop count is 27.

6.5.3 `pragma sst branch_predict` [float: C++ expression]

Similar to the way that loop counts can break the static analysis, if statements inside a loop skeletonized with `omp parallel` or `sst compute` can also be problematic. If the predicate depends on a variable declared *inside* the skeletonized block, the static analysis will break since it cannot predict when and how often to assume true or false. In contrast to the branch prediction pragma previously used, branch prediction pragmas inside a compute block must give a number between 0 and 1. This can either be a literal float or expression that computes a float value. Consider an example from CoMD:

```
1  #pragma sst branch_predict 0.2
2  if(r2 <= rCut2 && r2 > 0.0){
```

Inside the compute block, a compute may or may not occur depending on whether a particle distance is less than a cutoff. Based on the way CoMD constructs unit cells and halo regions, running CoMD shows that about 1 in 5 neighbor interactions are actually below the cutoff. Thus we given the branch prediction the hint 0.2.

Chapter 7

Issues and Limitations

7.1 Polling in applications

Use of testing non-blocking functions in a loop, such as:

```
1 while(!flag){  
2   MPI_Iprobe( 0, 0, MPI_COMM_WORLD, &flag, &status );  
3 }
```

For some configurations, virtual time never advances in the `MPI_Iprobe` call. This causes an infinite loop that never returns to the discrete event manager. Even if configured so that time progresses, the code will work but will take a very long time to run.

7.2 Fortran

SST/macro previously provided some experimental support for Fortran90 applications. This has been discontinued for the foreseeable future. For profiling existing apps written with Fortran, DUMPI traces can still be generated.

Chapter 8

Detailed Parameter Listings

The following chapter is organized by parameter namespaces. Tables in each namespace are organized as

Name (type)	Default if not given	Allowed Values	Description
-------------	----------------------	----------------	-------------

which lists the possible parameter names, allowed values, and brief descriptions. More detailed descriptions of particular parameter values are found in the documentation in previous chapters.

The allowed parameter types are:

int	Any integer
long	Any integer value, but guaranteed not to overflow for long integers
bool	Either “true” or “false” as lowercase string
time	Any valid float value followed by time units (s,ms,us,ns,ps)
freq	Any valid float value followed by frequency units (Hz, MHz, GHz)
bandwidth	Any valid float value followed by bandwidth units (b/s, B/s, Mb/s, MB/s, etc)
byte length	Any positive integer followed by length units (B, KB, MB, GB,TB)
string	An arbitrary string
vector of X	A vector of type X with entries separated by spaces
filepath	A valid filepath to an <i>existing</i> file, either absolute or relative
quantiy	A catch-all for a quantity with units. Any of frequency, bandwidth, byte length, or time can be given

8.1 Global namespace

sst_nthread (int)	1	Positive int	Only relevant for multi-threading. Specifying more threads than cores can lead to deadlock.
timestamp_resolution (time)	1ps		Specifies the length of time occupied by 1 timestamp tick - the smallest resolvable time difference. Numerical stability depends on this parameter matching the time scales of the simulation.
serialization_buffer_size (byte length)	512B		Size to allocate for buffers in a pool of serialization buffers. This should set be large enough to handle serialization of most messages, but not so large that significant space is wasted.
cpu_affinity (vector of int)	No default	Invalid cpu IDs give undefined behavior	When in multi-threading, specifies the list of core IDs that threads will be pinned to.

8.2 Namespace “topology”

geometry (vector of int)	No default	See Topology section	Geometry configuration of the topology. For details of the keyword, users should refer to Section 4
name (string)	No default	torus, dragonfly, fat_tree, butterfly, crossbar, tapered_fat_tree	The name of the topology to build. For details, see Section 4
seed (long)	System time		If no value given, random numbers for topology will be generated from system time
concentration (int)	1	Positive int	The number of nodes per network switch. For indirect networks, this is the number of nodes per leaf switch.
num_leaf_switches (int)	No default	Positive int	Only relevant for fat trees. This is the number of switches at the lowest level of the tree that are connected to compute nodes. Depending on how the fat tree is specified, this number may not be required.
k (int)	No default	int ≥ 2	The branching fraction of a fat tree. k=2 is a binary tree. k=4 is a quad-tree.
l (int)	No default	Positive int	The number of levels in a fat tree.
num_inj_switches_per_subtree	No default	Positive int	For a tapered tree, the number of injection switches, N_{inj} , within an aggregation tree that connect directly to compute nodes.
num_agg_switches_per_subtree	No default	Positive int	For a tapered tree, the number of aggregations witches per aggregation tree linking injection switches to the core.
num_agg_subtrees	No default	Positive int	For a tapered fat tree with 3 levels (injection, aggregation, core), this gives the number, N_{agg} , of aggregation subtrees. To find the total number, N_{tot} of injection (leaf) switches, we have $N_{tot} = N_{agg} \times N_{inj}$.
num_core_switches	No default	Positive int	The total number of core switches in a tapered tree linking the individual aggregation trees.
group_connections (int)	No default	Positive int	For dragonfly, the number of intergroup connections on each switch in a Dragonfly group
redundant (vector of int)	vector of 1's	Positive ints	For Cartesian topologies (hypercube, dragonfly, torus) this specifies a bandwidth (redundancy) multiplier for network links in each dimension.

8.3 Namespace “node”

model (string)	simple	simple	The type of node model (level of detail) for node-level operations
services (vector of strings)	Empty	Valid service names	For details, see section on distributed services in developer’s manual. Advanced feature.

8.3.1 Namespace “node.nic”

model (string)	No default	pisces, logP	The type of NIC model (level of detail) for modeling injection of messages (flows) to/from the network.
packetizer (string)	cut_through	merlin, simple, cut_through	The type of packetizer for injecting flows into the network. Merlin is part of sst-elements. Simple and cut-through use PISCES
packet_allocator (string)	pisces	pisces	The packet object used by the NIC. Certain type of statistics or congestion models may require custom packets with extra fields.
negligible_size (byte length)	256B		Messages (flows) smaller than size will not go through detailed congestion modeling. They will go through a simple analytic model to compute the delay.

Namespace “node.nic.delay_histogram”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
bin_size (quantity)	No default	Positive value	For the quantity being collected (bytes,time,etc), this determines how large a value range should be binned together.
num_bins (int)	20	Positive int	Determines the number of initial bins to allocate. If number is too small, simulation will allocate more. This is just a guess for optimization.
logarithmic (bool)	false		Whether to use a logarithmic axis for binning.

Namespace “node.nic.congestion_spyplot”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
type (string)	ascii	ascii, png	Whether to generate a simple traffic matrix in an ASCII file or to generate a PNG image of the traffic pattern
normalization (long)	Ignored	Positive int	Determines a normalization value that will effect the color scales of an output PNG file. By default, if not specified, the spyplot will determine the largest value and normalize all values for the output PNG to be 0.0-1.0. Values larger than 1.0 will have the same color as 1.0. Useful for having two PNG files with different values have the same color scales.

Namespace “node.nic.traffic_matrix”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
type (string)	ascii	ascii, png	Whether to generate a simple traffic matrix in an ASCII file or to generate a PNG image of the traffic pattern
normalization (long)	Ignored	Positive int	Determines a normalization value that will effect the color scales of an output PNG file. By default, if not specified, the spyplot will determine the largest value and normalize all values for the output PNG to be 0.0-1.0. Values larger than 1.0 will have the same color as 1.0. Useful for having two PNG files with different values have the same color scales.

Namespace “node.nic.local_bytes_sent”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
----------------------	------------	--	---

Namespace “node.nic.global_bytes_sent”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
----------------------	------------	--	---

Namespace “`node.nic.message_size_histogram`”

<code>fileroot</code> (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
<code>bin_size</code> (quantity)	No default	Positive value	For the quantity being collected (bytes,time,etc), this determines how large a value range should be binned together.
<code>num_bins</code> (int)	20	Positive int	Determines the number of initial bins to allocate. If number is too small, simulation will allocate more. This is just a guess for optimization.
<code>logarithmic</code> (bool)	false		Whether to use a logarithmic axis for binning.

Namespace “`node.nic.ejection`”

These parameters do not need to be specified, but can be given. Generally, the simulation assumes an infinite buffer size (unlimited memory) and no latency. All other parameters can be filled in from `node.nic.injection`.

stats (string)	null	delay_histogram, congestion_spyplot, multi, null	The type of statistics to collect from packets leaving or arriving at the NIC. Null indicates no statistics collected. For details on the other types of statistics, see Section 3.13
callbacks (vector of string)	No default	Any valid stats name	If stats is set to multi , then multiple different stats can be given here
arbitrator (string)	cut_through	null, simple, cut_through	Bandwidth arbitrator for PISCES congestion modeling. Null uses simple delays with no congestion. Simple uses store-and-forward that is cheap to compute, but can have severe latency errors for large packets. Cut-through approximates pipelining of flits across stages.
latency (time)	No default		If given, overwrites the send and credit latency parameters. Depending on component, the entire latency may be put on either the credits or the send.
bandwidth	No default		The bandwidth of the arbitrator
send_latency (time)	No default		The latency to send a packet to the next stage in the network. This can be omitted if the generic latency parameter is given (see above).
credit_latency (time)	No default		The latency to send a credit to the previous network stage. This can be omitted if the generic latency parameter is given (see above).
num_vc (int)	Computed	Positive int	If not specified, SST will estimate the number of virtual channels based on the topology and routing. If given and parameter is too small, the system can deadlock. If too large, buffer resources will be underutilized.
credits (byte length)	No default		The number of initial credits for the component. Corresponds to an input buffer on another component. In many cases, SST/macro can compute this from other parameters and fill in the value. In some cases, it will be required.
mtu (byte length)	1024B		The packet size. All messages (flows) will be broken into units of this size.

Namespace “node.nic.injection”

stats (string)	null	delay_histogram, null	The type of statistics to collect from packets leaving or arriving at the NIC. Null indicates no statistics collected. For details on the other types of statistics, see Section 3.13
arbitrator (string)	cut_through	null, simple, cut_through	Bandwidth arbitrator for PISCES congestion modeling. Null uses simple delays with no congestion. Simple uses store-and-forward that is cheap to compute, but can have severe latency errors for large packets. Cut-through approximates pipelining of flits across stages.
latency (time)	No default		If given, overwrites the send and credit latency parameters. Depending on component, the entire latency may be put on either the credits or the send.
bandwidth	No default		The bandwidth of the arbitrator
send_latency (time)	No default		The latency to send a packet to the next stage in the network. This can be omitted if the generic latency parameter is given (see above).
credit_latency (time)	No default		The latency to send a credit to the previous network stage. This can be omitted if the generic latency parameter is given (see above).
num_vc (int)	Computed	Positive int	If not specified, SST will estimate the number of virtual channels based on the topology and routing. If given and parameter is too small, the system can deadlock. If too large, buffer resources will be underutilized.
credits (byte length)	No default		The number of initial credits for the component. Corresponds to an input buffer on another component. In many cases, SST/macro can compute this from other parameters and fill in the value. In some cases, it will be required.
mtu (byte length)	1024B		The packet size. All messages (flows) will be broken into units of this size.

8.3.2 Namespace “node.memory”

model (string)	No default	logP, pisces	The type of memory model (level of detail) for modeling memory transactions.
arbitrator (string)	cut_through	null, simple, cut_through	The type of arbitrator. See arbitrator descriptions above.
latency (time)	No default		The latency of single memory operation
total_bandwidth	No default		The total memory bandwidth possible across all memory controllers.
max_single_bandwidth	Computed		The maximum memory bandwidth of a single stream of requests. Essentially the bandwidth of a single memory controller. If not given, this defaults the value of total_bandwidth.

8.3.3 Namespace “node.os”

compute_scheduler (string)	simple	simple, cpuset	The level of detail for scheduling compute tasks to cores. Simple looks for any empty core. cpuset allows bitmasks to be set for defining core affinities.
stack_size (byte length)	32KB		The size of user-space thread stack to allocate for each virtual application
stack_protect (bool)	false		Whether to put special mprotect regions around the thread stacks to turn stack overflow errors into seg faults
stack_chunk_size (byte length)	1MB		The size of memory to allocate at a time when allocating new thread stacks. Rather than allocating one thread stack at a time, multiple stacks are allocated and added to a pool as needed.

Namespace “node.os.call_graph”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
----------------------	------------	--	---

Namespace “node.os.ftq”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
epoch (time)	No default		The length of simulation time to treat as a single interval. This is much like a histogram bin, except the bin is partitioned into multiple categories. If too small, not enough data will be collected per interval to give reasonable looking results. If too large, significant changes in system activity over time will not be resolved.

8.3.4 Namespace “node.proc”

ncores (int)	No default	Positive int	The number of cores contained in a processor (socket). Total number of cores for a node is $ncores \times nsockets$.
frequency	No default		The baseline frequency of the node
parallelism (double)	1.0	Positive number	Fudge factor to account for superscalar processor. Number of flops per cycle performed by processor.

8.4 Namespace “netlink”

model (string)	null	null, pisces	The type of netlink congestion model to use. If null, no netlink blocks will be used.
concentration (int)	No default	Positive int	The number of nodes connected to a given netlink. The node concentration per switch must be greater than or equal to the node concentration per netlink.
num_tiles (int)	1	Positive int	The number of router tiles the netlink connects to. Only relevant for tiled router models. Advanced usage.

8.4.1 Namespace “netlink.injection”

arbitrator (string)	cut_through	null, simple, cut_through	Bandwidth arbitrator for PISCES congestion modeling. Null uses simple delays with no congestion. Simple uses store-and-forward that is cheap to compute, but can have severe latency errors for large packets. Cut-through approximates pipelining of flits across stages.
latency (time)	No default		If given, overwrites the send and credit latency parameters. Depending on component, the entire latency may be put on either the credits or the send.
bandwidth	No default		The bandwidth of the arbitrator
send_latency (time)	No default		The latency to send a packet to the next stage in the network. This can be omitted if the generic latency parameter is given (see above).
credit_latency (time)	No default		The latency to send a credit to the previous network stage. This can be omitted if the generic latency parameter is given (see above).
num_vc (int)	Computed	Positive int	If not specified, SST will estimate the number of virtual channels based on the topology and routing. If given and parameter is too small, the system can deadlock. If too large, buffer resources will be underutilized.
credits (byte length)	No default		The number of initial credits for the component. Corresponds to an input buffer on another component. In many cases, SST/macro can compute this from other parameters and fill in the value. In some cases, it will be required.
mtu (byte length)	1024B		The packet size. All messages (flows) will be broken into units of this size.

8.4.2 Namespace “netlink.ejection”

arbitrator (string)	cut_through	null, simple, cut_through	Bandwidth arbitrator for PISCES congestion modeling. Null uses simple delays with no congestion. Simple uses store-and-forward that is cheap to compute, but can have severe latency errors for large packets. Cut-through approximates pipelining of flits across stages.
latency (time)	No default		If given, overwrites the send and credit latency parameters. Depending on component, the entire latency may be put on either the credits or the send.
bandwidth	No default		The bandwidth of the arbitrator
send_latency (time)	No default		The latency to send a packet to the next stage in the network. This can be omitted if the generic latency parameter is given (see above).
credit_latency (time)	No default		The latency to send a credit to the previous network stage. This can be omitted if the generic latency parameter is given (see above).
num_vc (int)	Computed	Positive int	If not specified, SST will estimate the number of virtual channels based on the topology and routing. If given and parameter is too small, the system can deadlock. If too large, buffer resources will be underutilized.
credits (byte length)	No default		The number of initial credits for the component. Corresponds to an input buffer on another component. In many cases, SST/macro can compute this from other parameters and fill in the value. In some cases, it will be required.
mtu (byte length)	1024B		The packet size. All messages (flows) will be broken into units of this size.

8.5 Namespace “mpi”

test_delay (time)	0		The minimum time spent by MPI on each MPI_Test call
iprobe_delay (time)	0		The minimum time spent by MPI on each MPI_Iprobe call

8.5.1 Namespace “mpi.queue”

max_vshort_msg_size (byte length)	512B		The maximum size to use the very short message protocol. Small messages are sent eagerly using special pre-allocated mailbox buffers. Sends complete immediately.
max_eager_msg_size (byte length)	8KB		The maximum size to use the RDMA eager protocol. This also uses buffers to send message, but instead of using pre-allocated mailboxes, it coordinates an RDMA get. Sends complete immediately.
post_rdma_delay (time)	0		The minimum time spent by MPI posting each RDMA operation
post_header_delay (time)	0		The minimum time spent by MPI sending headers into pre-allocated mailboxes
poll_delay (time)	0		The minimum time spent by MPI each time it polls for incoming messages

8.6 Namespace “switch”

model (string)	No default	logP, pisces	The type of switch model (level of detail) for modeling network traffic.
buffer_size (byte length)	No default		The size of input and output buffers on each switch. This determines the number of credits available to other components

8.6.1 Namespace “switch.router”

name (string)	No default	minimal, valiant, ugal, fat_tree	The name of the routing algorithm to use for routing packets.
ugal_threshold (int)	0		The minimum number of network hops required before UGAL is considered. All path lengths less than value automatically use minimal.

8.6.2 Namespace “switch.output_buffer”

stats (string)	null	delay_histogram, byte_hops, null	The type of statistics to collect from packets leaving or arriving at the switch output buffers. Null indicates no statistics collected. For details on the other types of statistics, see Section 3.13
-------------------	------	-------------------------------------	---

Namespace “switch.output_buffer.delay_histogram”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
bin_size (quantity)	No default	Positive value	For the quantity being collected (bytes,time,etc), this determines how large a value range should be binned together.
num_bins (int)	20	Positive int	Determines the number of initial bins to allocate. If number is too small, simulation will allocate more. This is just a guess for optimization.
logarithmic (bool)	false		Whether to use a logarithmic axis for binning.

Namespace “switch.output_buffer.byte_hops”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
----------------------	------------	--	---

8.6.3 Namespace “switch.xbar”

stats (string)	null	bytes_sent, null	The type of statistics to collect from packets leaving or arriving at the switch crossbar. Null indicates no statistics collected. For details on the other types of statistics, see Section 3.13
arbitrator (string)	cut_through	null, simple, cut_through	Bandwidth arbitrator for PISCES congestion modeling. Null uses simple delays with no congestion. Simple uses store-and-forward that is cheap to compute, but can have severe latency errors for large packets. Cut-through approximates pipelining of flits across stages.
latency (time)	No default		If given, overwrites the send and credit latency parameters. Depending on component, the entire latency may be put on either the credits or the send.
bandwidth	No default		The bandwidth of the arbitrator
send_latency (time)	No default		The latency to send a packet to the next stage in the network. This can be omitted if the generic latency parameter is given (see above).
credit_latency (time)	No default		The latency to send a credit to the previous network stage. This can be omitted if the generic latency parameter is given (see above).
num_vc (int)	Computed	Positive int	If not specified, SST will estimate the number of virtual channels based on the topology and routing. If given and parameter is too small, the system can deadlock. If too large, buffer resources will be underutilized.
credits (byte length)	No default		The number of initial credits for the component. Corresponds to an input buffer on another component. In many cases, SST/macro can compute this from other parameters and fill in the value. In some cases, it will be required.
mtu (byte length)	1024B		The packet size. All messages (flows) will be broken into units of this size.

Namespace “switch.xbar.bytes_sent”

fileroot (string)	No default		The path of a file to be created (or overwritten) where results from the statistics will be placed.
----------------------	------------	--	---

8.6.4 Namespace “switch.link”

arbitrator (string)	cut_through	null, simple, cut_through	Bandwidth arbitrator for PISCES congestion modeling. Null uses simple delays with no congestion. Simple uses store-and-forward that is cheap to compute, but can have severe latency errors for large packets. Cut-through approximates pipelining of flits across stages.
latency (time)	No default		If given, overwrites the send and credit latency parameters. Depending on component, the entire latency may be put on either the credits or the send.
bandwidth	No default		The bandwidth of the arbitrator
send_latency (time)	No default		The latency to send a packet to the next stage in the network. This can be omitted if the generic latency parameter is given (see above).
credit_latency (time)	No default		The latency to send a credit to the previous network stage. This can be omitted if the generic latency parameter is given (see above).
num_vc (int)	Computed	Positive int	If not specified, SST will estimate the number of virtual channels based on the topology and routing. If given and parameter is too small, the system can deadlock. If too large, buffer resources will be underutilized.
credits (byte length)	No default		The number of initial credits for the component. Corresponds to an input buffer on another component. In many cases, SST/macro can compute this from other parameters and fill in the value. In some cases, it will be required.
mtu (byte length)	1024B		The packet size. All messages (flows) will be broken into units of this size.

8.6.5 Namespace “switch.ejection”

This namespace is not actually required. If unspecified, all of the values here will be filled in from `nic.injection`.

arbitrator (string)	cut_through	null, simple, cut_through	Bandwidth arbitrator for PISCES congestion modeling. Null uses simple delays with no congestion. Simple uses store-and-forward that is cheap to compute, but can have severe latency errors for large packets. Cut-through approximates pipelining of flits across stages.
latency (time)	No default		If given, overwrites the send and credit latency parameters. Depending on component, the entire latency may be put on either the credits or the send.
bandwidth	No default		The bandwidth of the arbitrator
send_latency (time)	No default		The latency to send a packet to the next stage in the network. This can be omitted if the generic latency parameter is given (see above).
credit_latency (time)	No default		The latency to send a credit to the previous network stage. This can be omitted if the generic latency parameter is given (see above).
num_vc (int)	Computed	Positive int	If not specified, SST will estimate the number of virtual channels based on the topology and routing. If given and parameter is too small, the system can deadlock. If too large, buffer resources will be underutilized.
credits (byte length)	No default		The number of initial credits for the component. Corresponds to an input buffer on another component. In many cases, SST/macro can compute this from other parameters and fill in the value. In some cases, it will be required.
mtu (byte length)	1024B		The packet size. All messages (flows) will be broken into units of this size.

8.7 Namespace “appN”

This is a series of namespaces `app1`, `app2`, and so on for each of the launched applications. These should be contained within the `node` namespace.

name (string)	No default	parsedumpi, cxx_full_main, cxx_empty_main	The name of the application to launch. Very few applications are built-in. Registration of external apps is shown starting in Section 3.5.
size (int)	No default	Positive int	The number of procs (MPI ranks) to launch. If launch_cmd given, this parameter is not required.
start (int)	0		The time at which a launch request for the application will be made
concentration (int)	1	Positive int	The number of procs (MPI ranks) per compute node
core_affinities (vector of int)	Empty		
launch_cmd (string)	No default	Valid aprun or srun	This uses a launch command as would be found with ALPS or SLURM launchers on real systems, e.g. aprun -n 4 -N 1
indexing (string)	block	block, random, cart, node_id, coordinate	The indexing scheme for assign proc ID (MPI rank number) to compute nodes
node_id_mapper_file (filepath)	No default		If using Node ID indexing, the file containing the node ID index list
random_indexer_seed (long)	System time		The seed to use for a random allocation. If not specified, system time is used.
allocation (string)	first_available	first_available, random, cart, node_id, coordinate	The scheme to use for allocating compute nodes to a given job.
random_allocation_seed (long)	System time		For random allocation policy. If unspecified, system time is used as the seed.
node_id_allocation_file (filepath)	No default		If using Node ID allocation, the file containing the list of node IDs to allocate for the job
dumpi_metaname (filepath)	No default		If running DUMPI trace, the location of the metafile for configuring trace replay
coordinate_file (filepath)	No default		If running using coordinate allocation or indexing, the path to the file containing the node coordinates of each proc (MPI rank)
cart_sizes (vector of int)	No default		Launch a contiguous block of nodes in a Cartesian topology. This gives the size of each dimension in the block.
cart_offsets (vector of int)	No default		Launch a contiguous block nodes in a Cartesian topology. This gives the offset in each dimension where the block begins.
parsedumpi_timescale (double)	1.0	Positive float	If running DUMPI traces, scale compute times by the given value. Values less than 1.0 speed up computation. Values greater than 1.0 slow down computation.
parsedumpi_terminate_percent (int)	100	1-100	Percent of trace. Can be used to terminate large traces early
host_compute_timer (bool)	False		Use the compute time on the host to estimate compute delays

otf2_metafile (string)	No default	string	The root file of an OTF2 trace.
otf2_timescale (double)	1.0	Positive float	If running OTF2 traces, scale compute times by the given value. Values less than 1.0 speed up computation. Values greater than 1.0 slow down computation.
otf2_print_mpi_calls (bool)	false		Print MPI calls found in the OTF2 trace
otf2_print_trace_events (bool)	false		Debugging flag that prints individual trace events (which includes details such as when an MPI call begins, ends, and when a collective begins and ends
otf2_print_time_deltas (bool)	false		Debugging flag that prints compute delays injected by the simulator
otf2_warn_unknown_callback (bool)	false		Debugging flag the prints unknown callbacks